

ECE469 - Introduction to ML

Midterm - Part 1

Chase A. Lotito - SIUC Undergraduate

October 21, 2024

Question 1. *Regression in machine learning.*

Solution.

(A)

```
HOUSING INFO
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

Figure 1: housing.info()

(B)

```
HOUSING DESCRIPTION
-----
      longitude    latitude  housing_median_age  ...  households  median_income  median_house_value
count  20640.000000  20640.000000    20640.000000  ...  20640.000000    20640.000000    20640.000000
mean   -119.569704    35.631861     28.639486  ...    499.539680      3.870671    206855.816909
std      2.003532      2.135952     12.585558  ...    382.329753      1.899822    115395.615874
min    -124.350000    32.540000      1.000000  ...      1.000000      0.499900     14999.000000
25%    -121.800000    33.930000     18.000000  ...    280.000000      2.563400    119600.000000
50%    -118.490000    34.260000     29.000000  ...    409.000000      3.534800    179700.000000
75%    -118.010000    37.710000     37.000000  ...    605.000000      4.743250    264725.000000
max    -114.310000    41.950000     52.000000  ...   6082.000000     15.000100    500001.000000

[8 rows x 9 columns]
```

Figure 2: housing.describe()

(C)

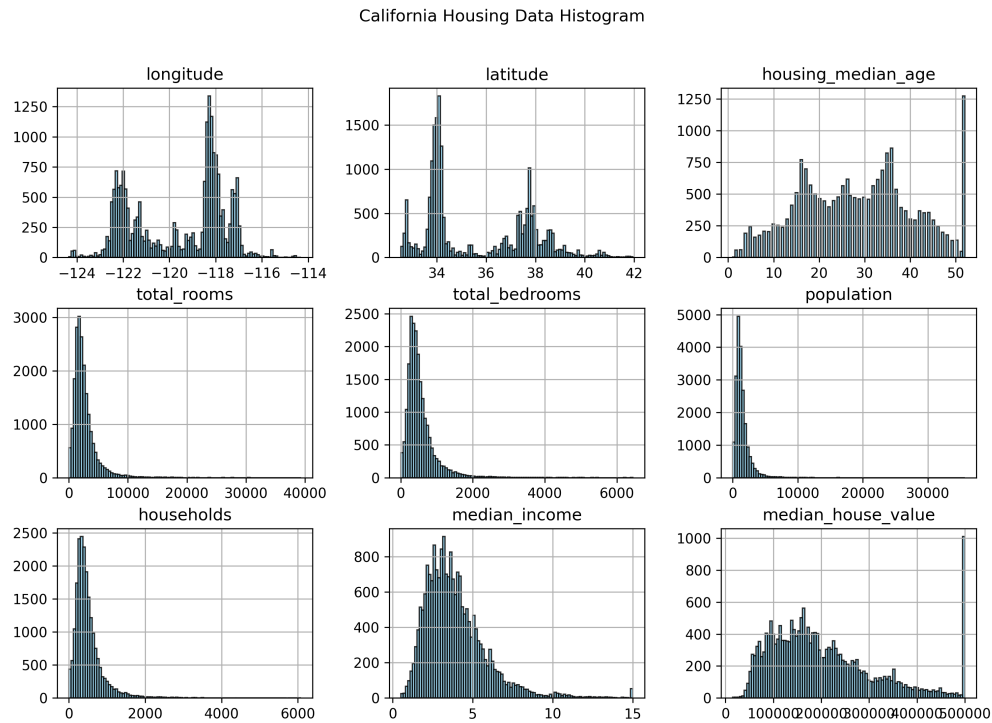


Figure 3: Housing Dataset Histogram

(D)

Using *SimpleImputer* from SciKit-Learn to replace missing data with the median value, and using *StandardScaler* from SciKit-Learn to standardize the dataset, the dataset looks as follows:

```
PREPROCESSED HOUSING DATA
-----
   people_per_house bedrooms_ratio rooms_per_house ... households median_income median_house_value
0          -0.049597       -1.029988      0.628559 ...    -0.977033         2.344766         2.129631
1          -0.092512       -0.888897      0.327041 ...     1.669961         2.332238         1.314156
2          -0.025843       -1.291686      1.155620 ...    -0.843637         1.782699         1.258693
3          -0.050329       -0.449613      0.156966 ...    -0.733781         0.932968         1.165100
4          -0.085616       -0.639087      0.344711 ...    -0.629157        -0.012881         1.172900
...          ...          ...          ...          ...          ...          ...
20635       -0.049110      0.165994      -0.155023 ...    -0.443449        -1.216128        -1.115804
20636       0.005021      0.021671      0.276881 ...    -1.008420        -0.691593        -1.124470
20637       -0.071735      0.021134      -0.090318 ...    -0.174042        -1.142593        -0.992746
20638       -0.091225      0.093467      -0.040211 ...    -0.393753        -1.054583        -1.058608
20639       -0.043682      0.113275      -0.070443 ...     0.079672        -0.780129        -1.017878

[20640 rows x 12 columns]
```

Figure 4: Preprocessed dataset

Below is the python code used for Question 1:

Listing 1: File: q1.py

```
1 """
2 Southern Illinois University Carbonele
3 Department of Electrical Engineering
4 -----
```

```

5 ECE469: Intro to Machine Learning
6 Midterm Exam: Question 1
7 Chase Lotito
8
9 10/14/2024
10
11 "California Housing Prices"
12 """
13
14 import pandas as pd
15 import numpy as np
16 import matplotlib.pyplot as plt
17
18 from pandas.core.common import random_state
19 from sklearn.impute import SimpleImputer
20 from sklearn.preprocessing import StandardScaler
21 from sklearn.preprocessing import PolynomialFeatures
22 from sklearn.model_selection import train_test_split
23 from sklearn.linear_model import LinearRegression
24 from sklearn.metrics import mean_squared_error
25
26 # import housing data from repository
27 url = "https://github.com/ageron/data/raw/main/housing/housing.csv"
28 housing = pd.read_csv(url)      # Store housing data in DataFrame
29 temp = housing
30
31 # important variables
32 plots_url = "./plots/q1/"
33
34 # (a) Use info() method to identify the attributes of this data-set
35 print('HOUSING INFO\n-----')
36 housing.info()
37
38
39 # (b) Use describe() method to identify and peek at a summary of
40 #     the numerical attributes.
41 housing_desc = housing.describe()
42 print('\nHOUSING DESCRIPTION\n-----')
43 print(housing_desc)
44
45
46 # (c) Use hist() method on the whole dataset and plot a histogram
47 #     for each numerical attribute. Notice that many histograms
48 #     are skewed right.
49 housing.hist(bins=100, color='skyblue', alpha=0.8, edgecolor='black', figsize
50             =(12,8))
51 plt.suptitle('California Housing Data Histogram')
52 #plt.savefig(plots_url + 'housing_histogram.png', dpi=300)
53 plt.close()
54
55 # (d) Clean and normalize/standardize the data-set to make it appropriate
56 #     for training a regression model. Creating training and test sets.
57 #     Create a copy of the data with only the numerical attributes by
58 #     excluding the text attribute ocean proximity from the data-set.
59
60 # remove ocean_proximity from the dataset
61 housing = housing.drop('ocean_proximity', axis=1)
62
63 # clean data via imputer; replacing missing values with median or mean

```

```

63 # this satisfies part (i) since total_bedrooms now is complete
64 simple_imputer = SimpleImputer(strategy='median')
65 housing_imputed = simple_imputer.fit_transform(housing)
66 housing = pd.DataFrame(housing_imputed, columns=housing.columns)
67
68 # STICKING PART H HERE SO I CAN ADD THESE BEFORE SPLITTING DATASET
69 # (h) Add three new attributes; (i) rooms per house = total rooms/households,
70 #     (ii) bedrooms ratio = total bedrooms/total rooms, and
71 #     (iii) people per house = population/households.
72
73 # assign housing DataFrame attributes to arrays
74 total_rooms = housing['total_rooms'].values
75 households = housing['households'].values
76 total_bedrooms = housing['total_bedrooms'].values
77 population = housing['population'].values
78
79 # calculate new attributes
80 rooms_per_house = total_rooms / households
81 bedrooms_ratio = total_bedrooms / total_rooms
82 people_per_house = population / households
83
84 # assign new attributes to housing DataFrame
85 # using .insert(0, ...) to stack each in front
86 housing.insert(0, 'rooms_per_house', rooms_per_house)
87 housing.insert(0, 'bedrooms_ratio', bedrooms_ratio)
88 housing.insert(0, 'people_per_house', people_per_house)
89
90 # standardize housing dataset
91 standard_scaler = StandardScaler()
92 housing_scaled = standard_scaler.fit_transform(housing)
93 housing = pd.DataFrame(housing_scaled, columns=housing.columns)
94 print('\nPREPROCESSED HOUSING DATA\n-----')
95 print(housing)
96
97 # extract input features and output target
98 x = housing[[
99     'people_per_house',
100    'bedrooms_ratio',
101    'rooms_per_house',
102    'longitude',
103    'latitude',
104    'housing_median_age',
105    'total_rooms',
106    'total_bedrooms',
107    'population',
108    'households',
109    'median_income'
110    ]].values
111 y = housing['median_house_value'].values
112
113 # split into training and testing datasets
114 test_ratio = 0.2
115 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=test_ratio,
116     random_state=42)
117
118 # (e) Because this data-set includes geographical information (latitude
119 #     and longitude), you are asked to create a scatterplot of all the
120 #     districts to visualize the geographical data in 2D space.

```

```

121 # get arrays containing latitude and longitudes
122 lat = housing[['latitude']].values
123 long = housing[['longitude']].values
124
125 # plot them as a scatterplot
126 plt.figure(figsize=(10,8))
127 plt.scatter(lat, long, marker='o', s=0.75, c='#32a852', alpha=0.95, label='2D Map'
128 )
129 plt.xlabel('Latitude')
130 plt.ylabel('Longitude')
131 plt.title(r'Normalized 2D Space Visualization of $housing.csv$')
132 plt.legend()
133 plt.gca().set_facecolor((0,0.1,0.8, 0.1))
134 #plt.savefig(plots_url + '2d_housing_scatter.png', dpi=300)
135 plt.close()
136
137 # (f) Compute the standard correlation coefficient between every pair of
138 # attributes using the corr() method.
139 corr_matrix = housing.corr()
140 corr_to_median_house_val = corr_matrix['median_house_value'].sort_values(ascending
141 =False)
142 print('\nCORRELATION of _____ TO median_house_value\n
143 -----')
144 print(corr_to_median_house_val)
145
146 # (g)
147 pd.plotting.scatter_matrix(housing, alpha=0.2, figsize=(20,16))
148 plt.suptitle('Scatter Matrix for $housing.csv$')
149 plt.tight_layout()
150 # plt.savefig(plots_url + 'housing_scatter_matrix.png', dpi=80)
151 plt.close()
152
153 # (i)
154
155 # initialize linear model
156 linear_regressor = LinearRegression()
157
158 # train linear model
159 linear_regressor.fit(x_train, y_train)
160
161 # test linear model
162 y_train_predicted = linear_regressor.predict(x_train)
163 y_test_predicted = linear_regressor.predict(x_test)
164
165
166 # now test for different dataset sizes
167 train_sizes = np.linspace(0.1, 1.0, 10) # Dataset sizes from 10% to 100% of the
168 training set
169 train_errors = []
170 test_errors = []
171
172 for train_size in train_sizes:
173     # Split data into training and test sets
174     X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
175 random_state=42)

```

```

175     # Use only a portion of the training set defined by `train_size`
176     X_train_subset = X_train[:int(train_size * len(X_train))]
177     y_train_subset = y_train[:int(train_size * len(y_train))]
178
179     # Train the model
180     lin_reg = LinearRegression()
181     lin_reg.fit(X_train_subset, y_train_subset)
182
183     # Make predictions
184     y_train_pred = lin_reg.predict(X_train_subset)
185     y_test_pred = lin_reg.predict(X_test)
186
187     # Calculate the errors
188     train_mse = mean_squared_error(y_train_subset, y_train_pred)
189     test_mse = mean_squared_error(y_test, y_test_pred)
190
191     # Store the errors
192     train_errors.append(train_mse)
193     test_errors.append(test_mse)
194
195 # PLOTTING RESULTS
196
197 plt.figure(figsize=(16,9))
198
199 # plot predicted against actual
200 ax1 = plt.subplot(1, 2, 1)
201 ax1.scatter(y_test, y_test_predicted, c='orange', marker='x', label='Learned Model',
202            alpha=0.6)
203 ax1.plot([min(y_test), max(y_test)], [min(y_test_predicted), max(y_test_predicted)],
204         c='blue', label='Ideal Model')
205 plt.xlabel('$\mathbf{y}_{test}$ Actual')
206 plt.ylabel('$\mathbf{y}_{test}$ Predicted')
207 plt.title('Learned Model Spread')
208 plt.legend()
209
210 ## Plot the Learned Weights Coefficients of the model
211 #weights = linear_regressor.coef_
212 #feature_names = ['people_per_house', 'bedrooms_ratio', 'rooms_per_house', '
213                 'longitude',
214                 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms',
215                 'population', 'households', 'median_income']
216
217 #ax2 = plt.subplot(1, 3, 2)
218 #ax2.barh(feature_names, weights)
219 #plt.xlabel("Model Weights $w_i$")
220 #plt.title("Linear Regression Feature Weights")
221
222 # Plot the learning curve
223 ax3 = plt.subplot(1, 2, 2)
224 ax3.plot(train_sizes, train_errors, label='Training Error', color='blue')
225 ax3.plot(train_sizes, test_errors, label='Test Error', color='green')
226 plt.title('Learning Curve')
227 plt.xlabel('Training Set Size (Ratio)')
228 plt.ylabel('Mean Squared Error')
229 plt.subplots_adjust(wspace=0.2)
230
231 plt.suptitle('Linear Model Results for $housing.csv$')
232 plt.legend()

```

```
230 plt.savefig(plots_url + "linear_model_results.png", dpi=300)
231 plt.close()
```

Question 2. *Classification in machine learning.*

Solution.

Below is the python code used for Question 2:

Listing 2: File: q2.py

```
1  """
2  Southern Illinois University Carbonale
3  Department of Electrical Engineering
4  -----
5  ECE469: Intro to Machine Learning
6  Midterm Exam: Question 2
7  Chase Lotito
8
9  10/15/2024
10
11  "MNIST"
12  """
13
14  import pandas as pd
15  import numpy as np
16  import matplotlib.pyplot as plt
17
18  from sklearn.datasets import fetch_openml
19  from sklearn.model_selection import train_test_split
20  from sklearn.preprocessing import StandardScaler
21  from sklearn.linear_model import LogisticRegression
22  from sklearn.neighbors import KNeighborsClassifier
23  from sklearn.metrics import root_mean_squared_error, accuracy_score, f1_score
24  from scipy.ndimage import shift
25
26  from sklearn.decomposition import PCA
27
28  # fetch the MNIST dataset
29  mnist = fetch_openml('mnist_784', as_frame=False)
30
31  # extract input features and output target
32  X = mnist['data']
33  y = mnist['target']
34
35  # the target values in y are strings, so we must
36  # first convert them to integers
37  y = y.astype(int)
38
39
40
41  # (b) Write a function that can shift an MNIST image
42  #     in any direction. Do this in all directions for
43  #     the training set, and append them to it.
44  def quad_direction_enricher(X: np.array, y: np.array, size: int, px: int):
45      """
46      To enrich an image dataset with 4 sets of
47      the original set shifted in all directions px.
48      (up, down, left, right)
49      """
50      temp = []
51      y_temp = []
```

```

52     for k in range(0, 4, 1):
53         for i in range(0, size, 1):
54             img = X[i].reshape(28,28)
55
56             if (k == 0):
57                 img = shift(img, (px, 0)) # shift up
58             elif (k == 1):
59                 img = shift(img, (-px, 0)) # shift down
60             elif (k == 2):
61                 img = shift(img, (0, px)) # shift right
62             elif (k == 3):
63                 img = shift(img, (0, -px)) # shift left
64             else:
65                 print('ERROR: image shift bounds error')
66
67             temp.append(img)
68             y_temp.append(y[i])
69
70     #enriched_X = np.array([img.flatten() for img in temp])
71     enriched_X = np.array([img.flatten() for img in temp])
72     enriched_Y = np.array(y_temp)
73     return enriched_X, enriched_Y
74
75 # enrich X
76 X, y = quad_direction_enricher(X, y, len(X), px=1)
77
78 print(f'len(X) = {len(X)}, len(y) = {len(y)}')
79
80 subset_size = 30000
81 # plot the examples of elements in the enriched dataset
82 #ax1 = plt.subplot(2, 4, 1)
83 #ax1.imshow(X[1].reshape(28,28))
84 #plt.title('Original NMIST')
85 #ax2 = plt.subplot(2, 4, 2)
86 #ax2.imshow(X[2].reshape(28,28))
87 #ax3 = plt.subplot(2, 4, 3)
88 #ax3.imshow(X[3].reshape(28,28))
89 #ax4 = plt.subplot(2, 4, 4)
90 #ax4.imshow(X[4].reshape(28,28))
91 #ax5 = plt.subplot(2, 4, 5)
92 #ax5.imshow(X[subset_size - 1].reshape(28,28))
93 #plt.title('Shifted NMIST', loc='center')
94 #ax6 = plt.subplot(2, 4, 6)
95 #ax6.imshow(X[2*subset_size - 2].reshape(28,28))
96 #ax7 = plt.subplot(2, 4, 7)
97 #ax7.imshow(X[3*subset_size - 3].reshape(28,28))
98 #ax8 = plt.subplot(2, 4, 8)
99 #ax8.imshow(X[4*subset_size - 4].reshape(28,28))
100 #plt.suptitle('Example Elements of Enriched MNIST')
101 #plt.show()
102 #plt.close()
103
104 # split dataset into training and testing sets
105 test_ratio = 0.2
106 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=test_ratio,
107     random_state=42)
108
109 # scale down size of dataset

```



```

110 test_size = int(subset_size*test_ratio)
111 train_size = int(subset_size*(1-test_ratio))
112 X_train = X_train[:train_size]
113 X_test = X_test[:test_size]
114 y_train = y_train[:train_size]
115 y_test = y_test[:test_size]
116
117 # next we want to standardize our data, but not
118 # targets, as to preserve y\in(0,9)
119 standard = StandardScaler()
120 X_train = np.array(X_train).reshape(len(X_train), -1)
121 X_test = np.array(X_test).reshape(len(X_test), -1)
122 X_train = standard.fit_transform(X_train)
123 X_test = standard.transform(X_test)
124
125 # train logistic regression classifier
126
127 pca = PCA(n_components=0.95)
128 X_train_pca = pca.fit_transform(X_train)
129 X_test_pca = pca.transform(X_test)
130 X_train = X_train_pca
131 X_test = X_test_pca
132
133 tolerance = 1e-3
134 classify = LogisticRegression(solver='lbfgs', penalty='l2', C=2, max_iter=1000,
135                               random_state=0)
136
137 # predict using model
138 y_test_pred = classify.predict(X_test)
139 y_train_pred = classify.predict(X_train)
140
141 # calc MSE
142 test_rmse = root_mean_squared_error(y_test, y_test_pred)
143 train_rmse = root_mean_squared_error(y_train, y_train_pred)
144
145 # determine model accuracy
146 test_accuracy = accuracy_score(y_test, y_test_pred) * 100
147 train_accuracy = accuracy_score(y_train, y_train_pred) * 100
148
149 # determine model f1 score
150 test_f1 = f1_score(y_test, y_test_pred, average='macro')
151 train_f1 = f1_score(y_train, y_train_pred, average='macro')
152
153 print('MNIST CLASSIFICATION REPORT')
154 print('#####')
155 print(f'Dataset Size : {subset_size}')
156 print(f'Enriched Size : {len(X)}')
157 print(f'Probabilistic Model Tolerance      : {tolerance*100}%')
158 print(f'Probabilistic Test      RMSE      : {test_rmse:.2f}')
159 print(f'Probabilistic Test  Accuracy      : {test_accuracy:.2f}%')
160 print(f'Probabilistic Test  F1 Score      : {test_f1:.2f}')
161 print(f'Probabilistic Train   RMSE      : {train_rmse:.2f}')
162 print(f'Probabilistic Train Accuracy      : {train_accuracy:.2f}%')
163 print(f'Probabilistic Train F1 Score      : {train_f1:.2f}')
164
165 # (c) KNN-based algorithms belong to the class of non-probabilistic classifiers.
166 #     You are asked to design a KNN-based classifier to classify handwritten
167 #     digits (0-9) in MNIST data-set.

```

```

168
169 # initialize the knn classifier
170 knn = KNeighborsClassifier(n_neighbors=5, metric='minkowski')
171
172 # train knn classifier
173 knn.fit(X_train, y_train)
174
175 # predict using knn classifier
176 y_test_pred_knn = knn.predict(X_test)
177 y_train_pred_knn = knn.predict(X_train)
178
179 # calc MSE
180 test_rmse_knn = root_mean_squared_error(y_test, y_test_pred_knn)
181 train_rmse_knn = root_mean_squared_error(y_train, y_train_pred_knn)
182
183 # determine model accuracy
184 test_accuracy_knn = accuracy_score(y_test, y_test_pred_knn) * 100
185 train_accuracy_knn = accuracy_score(y_train, y_train_pred_knn) * 100
186
187 # determine model f1 score
188 test_f1_knn = f1_score(y_test, y_test_pred_knn, average='macro')
189 train_f1_knn = f1_score(y_train, y_train_pred_knn, average='macro')
190
191 # add KNN model results to classification report
192 print(f'Non-probabilistic Test      RMSE   : {test_rmse_knn:.2f}')
```

```

193 print(f'Non-probabilistic Test  Accuracy : {test_accuracy_knn:.2f}%')
194 print(f'Non-probabilistic Test   F1 Score : {test_f1_knn:.2f}')
```

```

195 print(f'Non-probabilistic Train      RMSE   : {train_rmse_knn:.2f}')
```

```

196 print(f'Non-probabilistic Train Accuracy : {train_accuracy_knn:.2f}%')
```

```

197 print(f'Non-probabilistic Train F1 Score : {train_f1_knn:.2f}')
```