

ECE469 - Introduction to ML

Midterm - Part 1

Chase A. Lotito - SIUC Undergraduate

October 23, 2024

Question 1. *Regression in machine learning.*

Solution.

(A)

Unedited, the housing dataset contains 10 features, but eventually *ocean_proximity* will be dropped, and *median_house_value* will be chosen as the output target; all remaining columns will be the input features for the linear model.

```
HOUSING INFO
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms       20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

Figure 1: housing.info()

(B)

```
HOUSING DESCRIPTION
-----
count    longitude    latitude    housing_median_age    ...    households    median_income    median_house_value
mean     -119.569704      35.631861      28.639486    ...    499.539680      3.870671      206855.816909
std       2.003532        2.135952      12.585558    ...    382.329753      1.899822      115395.615874
min      -124.350000      32.540000       1.000000    ...     0.499900      0.499900      14999.000000
25%      -121.800000      33.930000      18.000000    ...    280.000000      2.563400      119600.000000
50%      -118.490000      34.260000      29.000000    ...    409.000000      3.534800      179700.000000
75%      -118.010000      37.710000      37.000000    ...    605.000000      4.743250      264725.000000
max      -114.310000      41.950000      52.000000    ...   6082.000000     15.000100     500001.000000

[8 rows x 9 columns]
```

Figure 2: housing.describe()

(C)

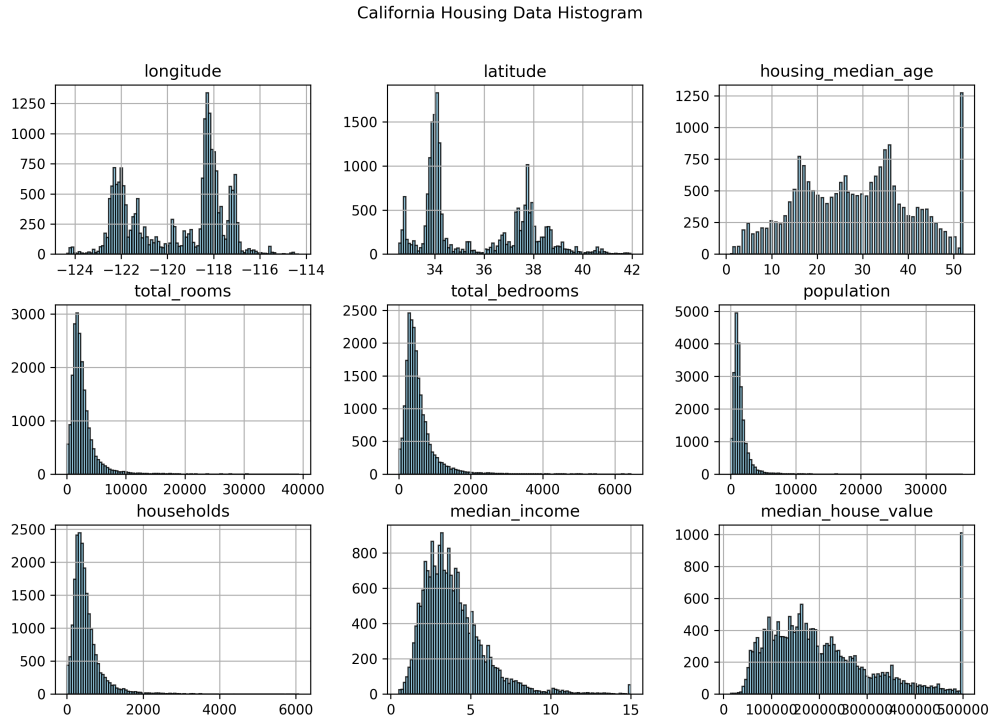


Figure 3: Housing Dataset Histogram

(D)

Using *SimpleImputer* from SciKit-Learn to replace missing data with the median value, and using *StandardScaler* from SciKit-Learn to standardize the dataset, the dataset looks as follows:

```
PREPROCESSED HOUSING DATA
-----
  people_per_house  bedrooms_ratio  rooms_per_house  ...  households  median_income  median_house_value
0          -0.049597        -1.029988         0.628559  ...    -0.977033         2.344766         2.129631
1          -0.092512        -0.888897         0.327041  ...     1.669961         2.332238         1.314156
2          -0.025843        -1.291686         1.155620  ...    -0.843637         1.782699         1.258693
3          -0.050329        -0.449613         0.156966  ...    -0.733781         0.932968         1.165100
4          -0.085616        -0.639087         0.344711  ...    -0.629157        -0.012881         1.172900
...          ...          ...          ...          ...          ...          ...
20635        -0.049110         0.165994        -0.155023  ...    -0.443449        -1.216128        -1.115804
20636         0.005021         0.021671         0.276881  ...    -1.008420        -0.691593        -1.124470
20637        -0.071735         0.021134        -0.090318  ...    -0.174042        -1.142593        -0.992746
20638        -0.091225         0.093467        -0.040211  ...    -0.393753        -1.054583        -1.058608
20639        -0.043682         0.113275        -0.070443  ...     0.079672        -0.780129        -1.017878

[20640 rows x 12 columns]
```

Figure 4: Preprocessed dataset

(E)

The standardized dataset when *latitude* and *longitude* are plotted against each other in a scatterplot shows a 2D map of all of the houses in the housing dataset.




Figure 5: 2D Housing Map - Latitude Longitude Scatterplot

(F)

Finding the correlation between the input features and *median_house_value* we find *median_income* has a strong positive influence on the value of a home, but for example, *total_bedrooms* is not as important.

```
CORRELATION of _____ TO median_house_value
-----
median_house_value    1.000000
median_income         0.688075
rooms_per_house       0.151948
total_rooms           0.134153
housing_median_age    0.105623
households            0.065843
total_bedrooms        0.049457
people_per_house      -0.023737
population            -0.024650
longitude             -0.045967
latitude              -0.144160
bedrooms_ratio        -0.233303
Name: median_house_value, dtype: float64
```

Figure 6: housing Scatter Matrix

(G)

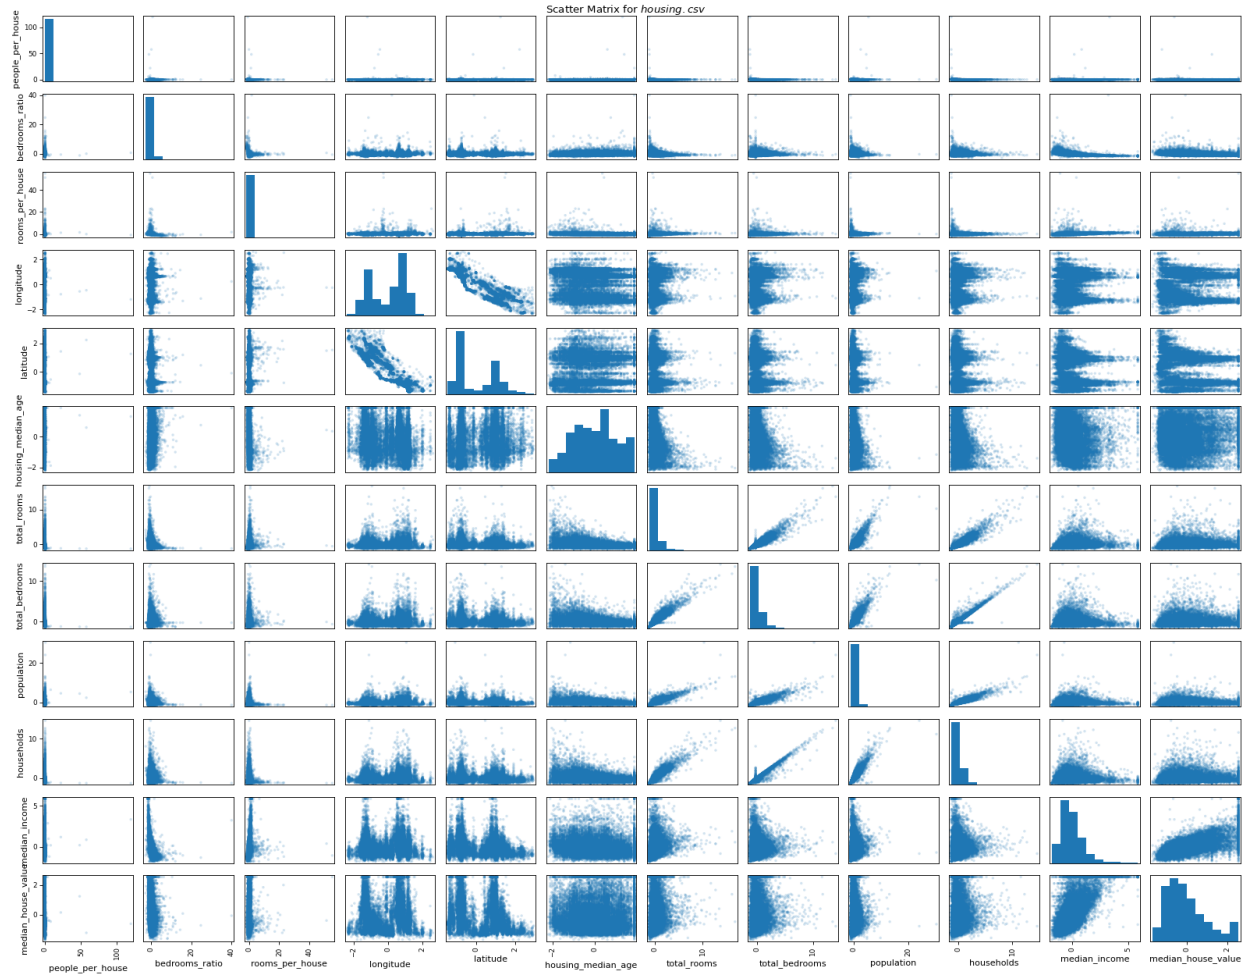


Figure 7: housing Scatter Matrix

(H)

The derived attributes *rooms per house*, *bedrooms ratio*, and *people per house* are created and inserted as input features in *q1.py* lines 73-88.

(I)

I chose to go the first route, where in lines 62-66, the missing values of *total_bedrooms*, and the rest of the dataset, were replaced with median values using *SimpleImputer*.

(J)

Below, in Fig. 8, we can see the predictions from our model plotted against the actual target values. Ideally, we would have all predictions equal to the actual values, and every orange marker would lie along the blue curve. However, the linear model struggles to predict accurately, and we find a spread.

Also, the *for* loop defined on line 171 trains the linear model for varying sizes of the training set size (almost a kind of cross-validation), where the mean-square-error is calculated for each size (10%-100%), and are plotted in Fig. 8. Here we see testing error falling as the model is trained on larger and larger datasets, but drops off beyond 80% as the model loses generalization.

It is worth noting that attempting to send the input features of the housing dataset into polynomial space, and then doing a linear regression on that dataset kills any kind of accuracy for any degree other than 1 (1 being nothing changed, still linear).

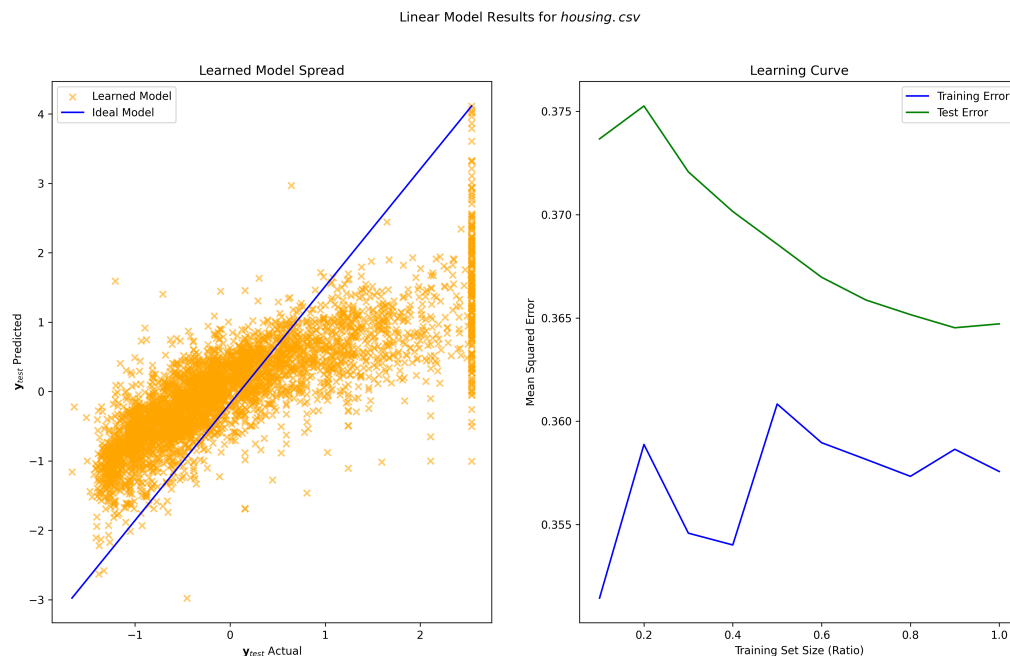


Figure 8: Linear Model Results: Model Spread (left), Learning Curve (right)

Below is the python code used for Question 1:

Listing 1: File: q1.py

```
1 """
2 Southern Illinois University Carbonale
3 Department of Electrical Engineering
4 -----
5 ECE469: Intro to Machine Learning
6 Midterm Exam: Question 1
```

```

7 Chase Lotito
8
9 10/14/2024
10
11 "California Housing Prices"
12 """
13
14 import pandas as pd
15 import numpy as np
16 import matplotlib.pyplot as plt
17
18 from pandas.core.common import random_state
19 from sklearn.impute import SimpleImputer
20 from sklearn.preprocessing import StandardScaler
21 from sklearn.preprocessing import PolynomialFeatures
22 from sklearn.model_selection import train_test_split
23 from sklearn.linear_model import LinearRegression
24 from sklearn.metrics import mean_squared_error
25
26 # import housing data from repository
27 url = "https://github.com/ageron/data/raw/main/housing/housing.csv"
28 housing = pd.read_csv(url) # Store housing data in DataFrame
29 temp = housing
30
31 # important variables
32 plots_url = "./plots/q1/"
33
34 # (a) Use info() method to identify the attributes of this data-set
35 print('HOUSING INFO\n-----')
36 housing.info()
37
38
39 # (b) Use describe() method to identify and peek at a summary of
40 # the numerical attributes.
41 housing_desc = housing.describe()
42 print('\nHOUSING DESCRIPTION\n-----')
43 print(housing_desc)
44
45
46 # (c) Use hist() method on the whole dataset and plot a histogram
47 # for each numerical attribute. Notice that many histograms
48 # are skewed right.
49 housing.hist(bins=100, color='skyblue', alpha=0.8, edgecolor='black', figsize
    =(12,8))
50 plt.suptitle('California Housing Data Histogram')
51 #plt.savefig(plots_url + 'housing_histogram.png', dpi=300)
52 plt.close()
53
54 # (d) Clean and normalize/standardize the data-set to make it appropriate
55 # for training a regression model. Creating training and test sets.
56 # Create a copy of the data with only the numerical attributes by
57 # excluding the text attribute ocean proximity from the data-set.
58
59 # remove ocean_proximity from the dataset
60 housing = housing.drop('ocean_proximity', axis=1)
61
62 # clean data via imputer; replacing missing values with median or mean
63 # this satisfies part (i) since total_bedrooms now is complete
64 simple_imputer = SimpleImputer(strategy='median')

```

```

65 housing_imputed = simple_imputer.fit_transform(housing)
66 housing = pd.DataFrame(housing_imputed, columns=housing.columns)
67
68 # STICKING PART H HERE SO I CAN ADD THESE BEFORE SPLITTING DATASET
69 # (h) Add three new attributes; (i) rooms per house = total rooms/households,
70 #     (ii) bedrooms ratio = total bedrooms/total rooms, and
71 #     (iii) people per house = population/households.
72
73 # assign housing DataFrame attributes to arrays
74 total_rooms = housing['total_rooms'].values
75 households = housing['households'].values
76 total_bedrooms = housing['total_bedrooms'].values
77 population = housing['population'].values
78
79 # calculate new attributes
80 rooms_per_house = total_rooms / households
81 bedrooms_ratio = total_bedrooms / total_rooms
82 people_per_house = population / households
83
84 # assign new attributes to housing DataFrame
85 # using .insert(0, ...) to stack each in front
86 housing.insert(0, 'rooms_per_house', rooms_per_house)
87 housing.insert(0, 'bedrooms_ratio', bedrooms_ratio)
88 housing.insert(0, 'people_per_house', people_per_house)
89
90 # standardize housing dataset
91 standard_scaler = StandardScaler()
92 housing_scaled = standard_scaler.fit_transform(housing)
93 housing = pd.DataFrame(housing_scaled, columns=housing.columns)
94 print('\nPREPROCESSED HOUSING DATA\n-----')
95 print(housing)
96
97 # extract input features and output target
98 x = housing[[
99     'people_per_house',
100     'bedrooms_ratio',
101     'rooms_per_house',
102     'longitude',
103     'latitude',
104     'housing_median_age',
105     'total_rooms',
106     'total_bedrooms',
107     'population',
108     'households',
109     'median_income'
110     ]].values
111 y = housing['median_house_value'].values
112
113 # split into training and testing datasets
114 test_ratio = 0.2
115 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=test_ratio,
116     random_state=42)
117
118 # (e) Because this data-set includes geographical information (latitude
119 #     and longitude), you are asked to create a scatterplot of all the
120 #     districts to visualize the geographical data in 2D space.
121
122 # get arrays containing latitude and longitudes
123 lat = housing[['latitude']].values

```

```

123 long = housing[['longitude']].values
124
125 # plot them as a scatterplot
126 plt.figure(figsize=(10,8))
127 plt.scatter(lat, long, marker='o', s=0.75, c='#32a852', alpha=0.95, label='2D Map'
128 )
129 plt.xlabel('Latitude')
130 plt.ylabel('Longitude')
131 plt.title(r'Normalized 2D Space Visualization of $housing.csv$')
132 plt.legend()
133 plt.gca().set_facecolor((0,0.1,0.8, 0.1))
134 #plt.savefig(plots_url + '2d_housing_scatter.png', dpi=300)
135 plt.close()
136
137 # (f) Compute the standard correlation coefficient between every pair of
138 #     attributes using the corr() method.
139 corr_matrix = housing.corr()
140 corr_to_median_house_val = corr_matrix['median_house_value'].sort_values(ascending
141 =False)
142 print('\nCORRELATION of _____ TO median_house_value\n
143 -----')
144 print(corr_to_median_house_val)
145
146 # (g)
147 pd.plotting.scatter_matrix(housing, alpha=0.2, figsize=(20,16))
148 plt.suptitle('Scatter Matrix for $housing.csv$')
149 plt.tight_layout()
150 # plt.savefig(plots_url + 'housing_scatter_matrix.png', dpi=80)
151 plt.close()
152
153 # (i)
154
155 # initialize linear model
156 linear_regressor = LinearRegression()
157
158 # train linear model
159 linear_regressor.fit(x_train, y_train)
160
161 # test linear model
162 y_train_predicted = linear_regressor.predict(x_train)
163 y_test_predicted = linear_regressor.predict(x_test)
164
165
166 # now test for different dataset sizes
167 train_sizes = np.linspace(0.1, 1.0, 10) # Dataset sizes from 10% to 100% of the
168 training set
169 train_errors = []
170 test_errors = []
171
172 for train_size in train_sizes:
173     # Split data into training and test sets
174     X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
175 random_state=42)
176
177     # Use only a portion of the training set defined by `train_size`
178     X_train_subset = X_train[:int(train_size * len(X_train))]

```



```

177     y_train_subset = y_train[:int(train_size * len(y_train))]
178
179     # Train the model
180     lin_reg = LinearRegression()
181     lin_reg.fit(X_train_subset, y_train_subset)
182
183     # Make predictions
184     y_train_pred = lin_reg.predict(X_train_subset)
185     y_test_pred = lin_reg.predict(X_test)
186
187     # Calculate the errors
188     train_mse = mean_squared_error(y_train_subset, y_train_pred)
189     test_mse = mean_squared_error(y_test, y_test_pred)
190
191     # Store the errors
192     train_errors.append(train_mse)
193     test_errors.append(test_mse)
194
195 # PLOTTING RESULTS
196
197 plt.figure(figsize=(16,9))
198
199 # plot predicted against actual
200 ax1 = plt.subplot(1, 2, 1)
201 ax1.scatter(y_test, y_test_predicted, c='orange', marker='x', label='Learned Model',
202            alpha=0.6)
203 ax1.plot([min(y_test), max(y_test)], [min(y_test_predicted), max(y_test_predicted)],
204         c='blue', label='Ideal Model')
205 plt.xlabel('$\mathbf{y}_{test}$ Actual')
206 plt.ylabel('$\mathbf{y}_{test}$ Predicted')
207 plt.title('Learned Model Spread')
208 plt.legend()
209
210 ## Plot the Learned Weights Coefficients of the model
211 #weights = linear_regressor.coef_
212 #feature_names = ['people_per_house', 'bedrooms_ratio', 'rooms_per_house', '
213     longitude',
214     'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms',
215     'population', 'households', 'median_income']
216
217 #ax2 = plt.subplot(1, 3, 2)
218 #ax2.barh(feature_names, weights)
219 #plt.xlabel("Model Weights $w_i$")
220 #plt.title("Linear Regression Feature Weights")
221
222 # Plot the learning curve
223 ax3 = plt.subplot(1, 2, 2)
224 ax3.plot(train_sizes, train_errors, label='Training Error', color='blue')
225 ax3.plot(train_sizes, test_errors, label='Test Error', color='green')
226 plt.title('Learning Curve')
227 plt.xlabel('Training Set Size (Ratio)')
228 plt.ylabel('Mean Squared Error')
229 plt.subplots_adjust(wspace=0.2)
230
231 plt.suptitle('Linear Model Results for $housing.csv$')
232 plt.legend()
233 plt.savefig(plots_url + "linear_model_results.png", dpi=300)
234 plt.close()

```

Question 2. *Classification in machine learning.*

Solution.

(A)

My best probabilistic classifier has 92.16% accuracy on the testing set, and my best non-probabilistic classifier has 97.12% on the testing set. Unfortunately, I was unable to achieve 97% accuracy for the probabilistic model on the testing set.

```
MNIST CLASSIFICATION REPORT
#####
Dataset Size : 70000
Probabilistic Model Tolerance      : 0.01%
Probabilistic Test      RMSE      : 1.16
Probabilistic Test  Accuracy      : 92.16%
Probabilistic Test  F1 Score      : 0.92
Probabilistic Train   RMSE      : 1.07
Probabilistic Train  Accuracy      : 93.27%
Probabilistic Train  F1 Score      : 0.93
Non-probabilistic Test      RMSE   : 0.76
Non-probabilistic Test  Accuracy   : 97.12%
Non-probabilistic Test  F1 Score   : 0.97
Non-probabilistic Train   RMSE     : 0.52
Non-probabilistic Train  Accuracy   : 98.58%
Non-probabilistic Train  F1 Score   : 0.99
```

Figure 9: MNIST Classification Results

(B)

I wrote a function called *quad_direction_enricher* (lines 43-73) which takes the MNIST dataset and shifts it a user-defined amount of pixels in each of the cardinal directions. The function returns an enriched version of the input features and the corresponding output target, which can then be appended onto the original dataset to increase the amount of data available to train. However, when trained over each of the now 350,000 datapoints, the model drops in accuracy from 92.16% to 89.97%. The reason for this I suspect is overfitting, or a reduction in the model's generalization.

```
MNIST CLASSIFICATION REPORT
#####
Dataset Size : 350000
Probabilistic Model Tolerance      : 0.01%
Probabilistic Test      RMSE      : 1.30
Probabilistic Test  Accuracy      : 89.97%
Probabilistic Test  F1 Score      : 0.90
Probabilistic Train   RMSE      : 1.28
Probabilistic Train  Accuracy      : 90.29%
Probabilistic Train  F1 Score      : 0.90
```

Figure 10: Expanded MNIST Classification Results

(C)

The non-probabilistic KNN classifier, using 3 neighbors, is in lines 130-155.

(D)

No matter what I did, the non-probabilistic KNN classifier always outperformed the probabilistic classifier; for my best model, the non-probabilistic classifier was 4.96% more accurate than the probabilistic classifier. In terms of computational complexity, the KNN classifier is much more computationally expensive and that is heavily dependent on how many neighbors you wish to use per datapoint, and the computational complexity exponentially increases for the increase in dataset size. For the 350,000 enriched dataset, my computer at 5000MHz could not finish the KNN classifier even after several minutes of waiting. So, there's a trade off. For critical applications that demand accuracy, using the non-probabilistic classifier is best, but for saving compute-power, it is best to choose the probabilistic model.




Figure 11: Probabilistic Classifier Confusion Matrix

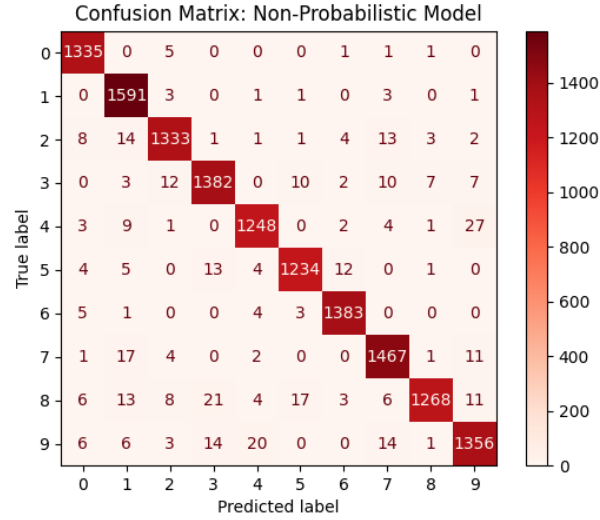


Figure 12: Non-probabilistic Classifier Confusion Matrix

Below is the python code used for Question 2:

Listing 2: File: q2.py

```
1  """
2  Southern Illinois University Carbondale
3  Department of Electrical Engineering
4  -----
5  ECE469: Intro to Machine Learning
6  Midterm Exam: Question 2
7  Chase Lotito
8
9  10/15/2024
10
11  "MNIST"
12  """
13
14  import pandas as pd
15  import numpy as np
16  import matplotlib.pyplot as plt
17
18  from sklearn.datasets import fetch_openml
19  from sklearn.model_selection import train_test_split, GridSearchCV
20  from sklearn.preprocessing import MinMaxScaler
21  from sklearn.linear_model import LogisticRegression
22  from sklearn.neighbors import KNeighborsClassifier
23  from sklearn.metrics import root_mean_squared_error, accuracy_score, f1_score
24  from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
25  from scipy.ndimage import shift
26
27  from sklearn.decomposition import PCA
28
29  # fetch the MNIST dataset
30  mnist = fetch_openml('mnist_784', as_frame=False)
31
32  # extract input features and output target
33  X = mnist['data']
34  y = mnist['target']
35
36  # the target values in y are strings, so we must
37  # first convert them to integers
38  y = y.astype(int)
39
40  # (b) Write a function that can shift an MNIST image
41  #     in any direction. Do this in all directions for
42  #     the training set, and append them to it.
43  def quad_direction_enricher(X: np.array, y: np.array, size: int, px: int):
44      """
45      To enrich an image dataset with 4 sets of
46      the original set shifted in all directions px.
47      (up, down, left, right)
48      """
49      enriched_X = []
50      y_temp = []
51      for k in range(0, 4, 1):
52          for i in range(0, size, 1):
53              img = X[i].reshape(28,28) # make image 28x28 matrix
54
55              # image shifting logic
56              if (k == 0):
```

```

57         img = shift(img, [px, 0]) # shift up
58     elif (k == 1):
59         img = shift(img, [-px, 0]) # shift down
60     elif (k == 2):
61         img = shift(img, [0, px]) # shift right
62     elif (k == 3):
63         img = shift(img, [0, -px]) # shift left
64     else:
65         print('ERROR: image shift bounds error')
66
67     img = img.flatten() # make image 1D array
68     enriched_X.append(img)
69     y_temp.append(y[i])
70
71     #enriched_X = np.array([img.flatten() for img in temp])
72     enriched_Y = np.array(y_temp)
73     return np.array(enriched_X), enriched_Y
74
75 # enrich X and y
76 e_X, e_y = quad_direction_enricher(X, y, len(X), px=1)
77
78 X = np.concatenate([X, e_X])
79 y = np.concatenate([y, e_y])
80
81
82 # scale down size of dataset
83 factor = 1
84 subset_size = 70000*factor
85 X = X[:subset_size]
86 y = y[:subset_size]
87
88 # split dataset into training and testing sets
89 test_ratio = 0.2
90 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=test_ratio,
91     random_state=42)
92
93 # next we want to standardize our data, but not
94 # targets, as to preserve y\in(0,9)
95 minmax = MinMaxScaler()
96 X_train = minmax.fit_transform(X_train)
97 X_test = minmax.transform(X_test)
98
99 # train logistic regression classifier
100
101 classify = LogisticRegression(multi_class='multinomial', solver='lbfgs', C=0.1,
102     max_iter=1000, random_state=42)
103 classify.fit(X_train, y_train)
104
105 # predict using model
106 y_test_pred = classify.predict(X_test)
107 y_train_pred = classify.predict(X_train)
108
109 # calc MSE
110 test_rmse = root_mean_squared_error(y_test, y_test_pred)
111 train_rmse = root_mean_squared_error(y_train, y_train_pred)
112
113 # determine model accuracy
114 test_accuracy = accuracy_score(y_test, y_test_pred) * 100
115 train_accuracy = accuracy_score(y_train, y_train_pred) * 100

```

```

114
115 # determine model f1 score
116 test_f1 = f1_score(y_test, y_test_pred, average='macro')
117 train_f1 = f1_score(y_train, y_train_pred, average='macro')
118
119 print('MNIST CLASSIFICATION REPORT')
120 print('#####')
121 print(f'Dataset Size : {subset_size}')
122 print(f'Probabilistic Model Tolerance      : {0.01}%')
123 print(f'Probabilistic Test      RMSE      : {test_rmse:.2f}%')
124 print(f'Probabilistic Test      Accuracy   : {test_accuracy:.2f}%')
125 print(f'Probabilistic Test      F1 Score    : {test_f1:.2f}%')
126 print(f'Probabilistic Train      RMSE      : {train_rmse:.2f}%')
127 print(f'Probabilistic Train      Accuracy   : {train_accuracy:.2f}%')
128 print(f'Probabilistic Train      F1 Score    : {train_f1:.2f}%')
129
130 # (c) KNN-based algorithms belong to the class of non-probabilistic classifiers.
131 #     You are asked to design a KNN-based classifier to classify handwritten
132 #     digits (0-9) in MNIST data-set.
133
134 # initialize the knn classifier
135 n=3
136 knn = KNeighborsClassifier(n_neighbors=n, metric='minkowski')
137
138 # train knn classifier
139 knn.fit(X_train, y_train)
140
141 # predict using knn classifier
142 y_test_pred_knn = knn.predict(X_test)
143 y_train_pred_knn = knn.predict(X_train)
144
145 # calc MSE
146 test_rmse_knn = root_mean_squared_error(y_test, y_test_pred_knn)
147 train_rmse_knn = root_mean_squared_error(y_train, y_train_pred_knn)
148
149 # determine model accuracy
150 test_accuracy_knn = accuracy_score(y_test, y_test_pred_knn) * 100
151 train_accuracy_knn = accuracy_score(y_train, y_train_pred_knn) * 100
152
153 # determine model f1 score
154 test_f1_knn = f1_score(y_test, y_test_pred_knn, average='macro')
155 train_f1_knn = f1_score(y_train, y_train_pred_knn, average='macro')
156
157 # add KNN model results to finish classification report
158 print(f'Non-probabilistic Test      RMSE      : {test_rmse_knn:.2f}%')
159 print(f'Non-probabilistic Test      Accuracy   : {test_accuracy_knn:.2f}%')
160 print(f'Non-probabilistic Test      F1 Score    : {test_f1_knn:.2f}%')
161 print(f'Non-probabilistic Train      RMSE      : {train_rmse_knn:.2f}%')
162 print(f'Non-probabilistic Train      Accuracy   : {train_accuracy_knn:.2f}%')
163 print(f'Non-probabilistic Train      F1 Score    : {train_f1_knn:.2f}%')
164
165
166 # (D) Confusion Matrix
167 prob_cmatrix = confusion_matrix(y_pred=y_test_pred, y_true=y_test)
168 ConfusionMatrixDisplay(confusion_matrix=prob_cmatrix).plot(cmap='Blues')
169 plt.title('Confusion Matrix: Probabilistic Model')
170 plt.show()
171 plt.close()
172

```

```
173 nonprob_cmatrix = confusion_matrix(y_pred=y_test_pred_knn, y_true=y_test)
174 ConfusionMatrixDisplay(confusion_matrix=nonprob_cmatrix).plot(cmap='Reds')
175 plt.title('Confusion Matrix: Non-Probabilistic Model')
176 plt.show()
177 plt.close()
```