

Question 1: Select ALL correct choices.

[1.1] Ans: **B**

[1.2] Ans: **A, B**

[1.3] Ans: **A, C**

[1.4] Ans: **A, B, D**

[1.5] Ans: **A, B, C**

Question 2: A linear ML model can be written as:

$$f(\mathbf{x}, \mathbf{w}) = \sum_{i=0}^n w_i x_i = \mathbf{w}^T \mathbf{x}$$

The loss function can be written as:

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^n \left[f(\mathbf{x}^{(i)}, \mathbf{w}) - y^{(i)} \right]^2$$

[2.1] Show analytically that the optimal weight vector that minimizes the cost function $J(\mathbf{w})$ is:

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Solution.

Since our model is linear, we can write the cost function as:

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^n \left[\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)} \right]^2 \quad (1)$$

The product $\mathbf{w}^T \mathbf{x}^{(i)} = \mathbf{X} \mathbf{w}$, $\forall i \in \{1, 2, \dots, n\}$ since the LHS implies the matrix multiplication of the RHS, as \mathbf{X} is the matrix of all input entries $\mathbf{x}^{(i)}$. So, Eq. 1 can be written as:

$$J(\mathbf{w}) = \frac{1}{m} [\mathbf{X} \mathbf{w} - \mathbf{y}]^2 \quad (2)$$

The inside of the brackets is just a vector, and the square of a vector is the norm of a vector, so we can reduce Eq. 2:

$$J(\mathbf{w}) = \frac{1}{m} \|\mathbf{X} \mathbf{w} - \mathbf{y}\| \quad (3)$$

$$= \frac{1}{m} (\mathbf{X} \mathbf{w} - \mathbf{y})^T (\mathbf{X} \mathbf{w} - \mathbf{y}) \quad (4)$$

$$= \frac{1}{m} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}) \quad (5)$$

Now to optimize the cost with respect to weights, we can take the gradient of $J(\mathbf{w})$ w.r.t. the weights $\mathbf{w}^{(i)}$.

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \frac{1}{m} \nabla_{\mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}) \quad (6)$$

$$= \frac{1}{m} \nabla_{\mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w}) \quad (7)$$

$$= \frac{1}{m} [\nabla_{\mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}) - \nabla_{\mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{y}) - \nabla_{\mathbf{w}} (\mathbf{y}^T \mathbf{X} \mathbf{w})] \quad (8)$$

The gradients $\nabla_{\mathbf{w}}$ are simply derivatives of each matrix function w.r.t. \mathbf{w} , which can be computed using equations (69) and (81) from *The Matrix Cookbook* [2].

$$= \frac{1}{m} \left[\frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}) - \frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{y}) - \frac{\partial}{\partial \mathbf{w}} (\mathbf{y}^T \mathbf{X} \mathbf{w}) \right] \quad (9)$$

$$= \frac{1}{m} ((\mathbf{X}^T \mathbf{X} + (\mathbf{X}^T \mathbf{X})^T) \mathbf{w} - \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X}) \quad (10)$$

$$= \frac{1}{m} ((\mathbf{X}^T \mathbf{X} + \mathbf{X}^T (\mathbf{X}^T)^T) \mathbf{w} - 2\mathbf{X}^T \mathbf{y}) \quad (11)$$

$$= \frac{1}{m} ((\mathbf{X}^T \mathbf{X} + \mathbf{X}^T \mathbf{X}) \mathbf{w} - 2\mathbf{X}^T \mathbf{y}) \quad (12)$$

$$= \frac{1}{m} (2\mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{X}^T \mathbf{y}) \quad (13)$$

$$= \frac{2}{m} (\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y}) \quad (14)$$

We find the minimum when $\nabla_{\mathbf{w}} J(\mathbf{w}) = 0$,

$$\frac{2}{m} (\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y}) = 0 \quad (15)$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y} = 0 \quad (16)$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y} \quad (17)$$

$$(\mathbf{X}^T \mathbf{X})^{-1} (\mathbf{X}^T \mathbf{X} \mathbf{w}) = (\mathbf{X}^T \mathbf{X})^{-1} (\mathbf{X}^T \mathbf{y}) \quad (18)$$

$$\mathbf{I} \mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad (19)$$

$$\boxed{\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{w}^*} \quad (20)$$

Without advanced methods, multiplying two $(n \times n)$ matrices is of $O(n^3)$ complexity [1]. For very large data sets, computing \mathbf{w} would be enormously computationally expensive.

[2.2] Develop a pseudo-codes for implementing the batch gradient descent, stochastic gradient descent, and mini-batch gradient descent algorithms to train the above linear model.

Solution.

Batch Gradient Descent

```

1  a = learning rate
2  N = size of dataset
3
4  # Psuedocode for Batch Gradient Descent
5  REPEAT UNTIL CONVERGENCE
6  {
7  for all j:
8      temp_j = w_j - (a / N) * sum( i=1..N, gradient_w(J(
          x_i, w_j) )
9
10 w_j = temp_j
11 }

```

Stochastic Gradient Descent

```

1  a = learning rate
2
3  # Pseudocode for Stochastic Gradient Descent
4  REPEAT UNTIL CONVERGENCE
5  {
6  for all i:
7      w_j = w_j - a * gradient_w(J(x_i, w_j))
8  }

```

Mini-Batch Gradient Descent

```

1  a = learning rate
2  b = batch size
3  m = rows in training set
4
5  # Psuedocode for Mini-Batch Gradient Descent
6  REPEAT UNTIL CONVERGENCE
7  {
8  # i indexes will be multiples of batch size
9  for i = (1*b + 1, 2*b + 1, ... , (m-1)*b + 1):
10      w_j = w_j - (a / b) * sum(k=i..i+(b-1), gradient_w(J(
          x_k, w_j)))
11 }

```

[2.3] Discuss the performance versus computational complexity of each of the above algorithms.

Solution.

Batch-Gradient Descent (BDG) is the most computationally complex since the algorithm has to sum over the entire dataset *every single update* of w_j . Stochastic Gradient Descent (SDG) is much faster than BDG since each update of w_j only looks over a single data row in the X matrix, instead of all of them. Mini-Batch Gradient Descent (MBDG) is inbetween BDG and SDG in terms of computational

complexity, since MBDG updates w_j after summing the gradients of a specified batch size $b \in (1, \dots, N)$.

A ranking of each algorithm from least computationally complex to least is as follows:

- (1) Stochastic
- (2) Mini-Batch
- (3) Batch

Question 3: Housing data preprocessing.

Solution.

```
1  # Chase Lotito - SIUC Fall 2024 - ECE469: Intro to
    Machine Learning
2  # HW1 - Question 3
3
4  import numpy as np
5  import pandas as pd
6  from sklearn.preprocessing import OrdinalEncoder      # For
    encoding categorical features
7  from sklearn.impute import SimpleImputer             # For
    adding missing values
8  from sklearn.preprocessing import StandardScaler     # For
    standardizing data
9
10 # (A) Get housing data
11 RAW_DATA = 'https://github.com/ageron/data/raw/main/
    housing/housing.csv'
12 housing = pd.read_csv(RAW_DATA)
13
14 # (B) Choose input features and output features (saved
    into numpy.ndarray type)
15 X = housing[
16     ['longitude',
17     'latitude',
18     'housing_median_age',
19     'total_rooms',
20     'total_bedrooms',
21     'population',
22     'households',
23     'median_income',
24     'ocean_proximity']]
25     ].values
26 Y = housing[['median_house_value']].values
27
28 # (C) Ocean Proximity is a categorical feature. Drop it
    or transform into numerical values (encode).
29
```

```
30 # Isolate the ocean_proximity data in input data X
31 ocean_proximity = X[:,8].reshape(-1,1) # reshape(-1,1)
    to make 2D array for Ordinal
32 # Initialize the ordinal encoder
33 ordinal_encoder = OrdinalEncoder()
34 # Encode the ocean_proximity strings into numerical data
35 encoded_ocean = ordinal_encoder.fit_transform(
    ocean_proximity)
36 # Put the encoded version of ocean_proximity into input
    data X
37 X[:,8] = encoded_ocean.flatten() # flatten to add
    1D version of array back into X
38
39 # (D) Clean the data by either dropping or replacing
    missing values
40
41 # Initialized SimpleImputer, will use the median to add
    missing entries
42 simple_imputer = SimpleImputer(strategy='median')
43
44 # Change X np ndarray into a Pandas Dataframe to use
    SimpleImputer
45 dX = pd.DataFrame(X)
46 dY = pd.DataFrame(Y)
47
48 # Perform SimpleImputer transformation, for both inputs X
    and outputs Y
49 imputed_data = simple_imputer.fit_transform(dX)
50 X = imputed_data
51 imputed_data = simple_imputer.fit_transform(dY)
52 Y = imputed_data
53
54 # (E) Carry out feature scaling either via normalization
    or standardization.
55 std_scaler = StandardScaler()
56 scaled_data = std_scaler.fit_transform(X)
57 X = scaled_data
58 scaled_data = std_scaler.fit_transform(Y)
59 Y = scaled_data
60
61 # (F) Create a training dataset and testing dataset
62 def shuffle_and_split_data(data, test_ratio):
63     shuffled_indices = np.random.permutation(len(data))
64     test_set_size = int(len(data) * test_ratio)
65     test_indices = shuffled_indices[:test_set_size]
66     train_indices = shuffled_indices[test_set_size:]
67     return data.iloc[train_indices], data.iloc[
        test_indices]
```

```
68
69 # first, recombine X and Y into augmented matrix
70 housing_data = np.hstack((X,Y))
71
72 # split into testing and training set (both outputted as
   pd.DataFrames)
73 housing_training, housing_testing =
   shuffle_and_split_data(pd.DataFrame(housing_data),
   0.2)
74 print('TRAINING:')
75 print(housing_training)
76 print('TESTING:')
77 print(housing_testing)
```

References

- [1] Andy He and Evan Williams. Computational complexity of matrix multiplication. <https://www.cs.cornell.edu/courses/cs6810/2023fa/Matrix.pdf>, Fall 2023. Accessed: 2024-09-10.
- [2] Kaare Brandt Petersen and Michael Syskind Pederson. The matrix cookbook. Distributed by University of Waterloo, November 2012.