

ECE469 Homework 2

Chase A. Lotito

October 5, 2024

Question 1

Solution.

q1.py

```
1 #####
2 # Chase Lotito - SIUC F24 #
3 # ECE469 - Intro to ML    #
4 # HW2 - Question 1       #
5 #####
6
7 # IMPORT LIBRARIES
8 import numpy as np
9 import pandas as pd
10 import matplotlib.pyplot as plt
11 from sklearn.preprocessing import PolynomialFeatures
12 from sklearn.model_selection import cross_val_score,
    train_test_split, KFold
13 from sklearn.linear_model import LinearRegression
14 from sklearn.linear_model import Ridge
15 from sklearn.metrics import mean_squared_error
16
17 # Read in provided csv data to pandas dataframe
18 RAW_DATA_PATH = 'C:/Users/cloti/OneDrive/Desktop/CODE/datasets/
    datasetHW2P1.csv'
19 data = pd.read_csv(RAW_DATA_PATH)
20
21 test_size = 0.2
22
23
24 # (A) SPLIT DATASET TO CREATE TWO SUB DATASETS FOR TRAINING AND
    TESTING
25 x = data['x'].values
26 y = data['y'].values
27
28 x = x.reshape(-1, 1)
```

```

29
30 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=
    test_size, random_state=42)
31
32
33
34
35 # (B) USE POLYNOMIAL REGRESSION TO FIT 5 POLYNOMIAL MODELS (DEG. 1 -
    DEG. 5)
36 deg = [1, 2, 3, 4, 5]
37 mse_train_arr = []
38 mse_test_arr = []
39
40 for i in deg:
41     # transform input features into polynomial features
42     poly = PolynomialFeatures(degree=i) # initialize polynomial
43     x_train_poly = poly.fit_transform(x_train)
44     x_test_poly = poly.transform(x_test)
45
46     # fit model
47     poly_regressor = LinearRegression()
48     poly_regressor.fit(x_train_poly, y_train)
49
50     # test
51     y_train_predicted = poly_regressor.predict(x_train_poly)
52     y_test_predicted = poly_regressor.predict(x_test_poly)
53
54     # calc mse
55     mse_train = mean_squared_error(y_train, y_train_predicted)
56     mse_test = mean_squared_error(y_test, y_test_predicted)
57
58     # add mse's to arrays for later plotting
59     mse_train_arr.append(mse_train)
60     mse_test_arr.append(mse_test)
61
62     # print mse for training and testing
63     print(f"Degree: {i}")
64     print(f"MSE Training: {mse_train}")
65     print(f"MSE Testing: {mse_test}")
66
67 # plotting mse for training and testing against model complexity
68 plt.plot(deg, mse_train_arr, label='Training')
69 plt.plot(deg, mse_test_arr, label='Testing')
70 plt.xlabel('Model Complexity (Degree)')
71 plt.ylabel('Mean Square Error')
72 plt.title('Model Error v. Model Complexity')
73 plt.legend()
74 #plt.show()

```

```

75 plt.close()
76
77
78
79
80 # (C) USE 10-FOLD CROSS-VALIDATION TO FIND THE MODEL WHICH OPTIMALLY
    FITS GIVEN DATASET.
81 #     PLOT THE TRAINING, CROSS-VALIDATION, AND TESTING ERRORS
    AGAINST MODEL COMPLEXITY.
82
83 # parameters
84 k = 10 # for 10-fold cross-validation
85 best_degree = 1 # track best degree
86 max_degree = 5
87 best_mse = float('inf')
88 mse_per_degree = []
89
90 # perform k-fold cross-validation for each polynomial degree
91 for degree in range(1, max_degree + 1):
92     kf = KFold(n_splits=k, shuffle=True, random_state=42)
93     mse_fold_values = []
94     k_iter = 1
95
96     # initialize a plot for regression fit
97     all_y_pred = np.zeros(x.shape)
98
99     # Loop through each fold
100    for train_index, test_index in kf.split(x):
101        x_train, x_test = x[train_index], x[test_index]
102        y_train, y_test = y[train_index], y[test_index]
103
104        # transform original x data into polynomial features for
            current degree
105        poly = PolynomialFeatures(degree=degree) # here degree is
            the parent loop iterator
106        x_train_poly = poly.fit_transform(x_train)
107        x_test_poly = poly.transform(x_test)
108
109        # initialize and fit linear regression on polynomial space
            input features
110        poly_regressor = LinearRegression()
111        poly_regressor.fit(x_train_poly, y_train) # <--
            TRAINING MODEL HERE
112
113        # predict using trained model
114        y_test_pred = poly_regressor.predict(x_test_poly)
115
116        # calc mse for test set

```

```

117         mse_test = mean_squared_error(y_test, y_test_pred)
118         mse_fold_values.append(mse_test)
119
120     # calculate avg mse for all folds for current deg
121     avg_mse = np.mean(mse_fold_values)
122     mse_per_degree.append(avg_mse)
123
124     # check if current deg has lowest avg mse
125     if avg_mse < best_mse:
126         best_mse = avg_mse
127         best_degree = degree
128
129
130     # Generate new plot for degree
131     plt.figure(figsize=(12,8))
132     plt.scatter(x,y,s=10,color='blue', label='Original Data', alpha
133                =0.6)
134
135     # plot polynomial fit for degree
136     sorted_x = np.sort(x, axis=0)
137     sorted_x_poly = poly.transform(sorted_x)
138     plt.plot(sorted_x, poly_regressor.predict(sorted_x_poly), color=
139              'red', label=f"Degree {degree}", alpha = 0.7)
140
141     # finalize plot for degree
142     plt.title(f"Polynomial Regression (Degree {degree})")
143     plt.xlabel('x')
144     plt.ylabel('y')
145     plt.legend(loc='best')
146     plt.grid(True)
147     plt.savefig(f"poly_reg_deg{degree}_fold{k_iter}.png", dpi=300)
148     plt.close()
149
150     # print best degree and mse
151     print(f"Best Polynomial Degree: {best_degree}")
152
153     # plot the mse v. polynomial degree
154     plt.figure(figsize=(10,6))
155     plt.plot([1,2,3,4,5], mse_per_degree, marker='o', color='b', label='
156             Cross-Validation MSE', alpha=0.75)
157     plt.plot(deg, mse_train_arr, label='Training MSE', marker='+', alpha
158            =0.75)
159     plt.plot(deg, mse_test_arr, label='Testing MSE', marker='x', alpha
160            =0.75)
161     plt.title("MSE vs Model Complexity")
162     plt.xlabel('Polynomial Degree')
163     plt.ylabel('Mean Squared Error (MSE)')
164     plt.xticks(range(1, max_degree + 1))

```

```

160 plt.legend()
161 plt.grid(True)
162 plt.savefig('plots\\mse_vs_complexity.png', dpi=300)
163 plt.close()
164
165
166
167
168 # (D) CONSIDER A 4-DEGREE POLYNOMIAL AS YOUR MODEL.
169 #     USE RIDGE REGRESSION AND FIND BEST HYPERPARAMETER
170 #     LAMBDA VIA 10-FOLD CROSS-VALIDATION. PLOT THE CROSS
171 #     VALIDATION ERROR VERSUS LN(LAMBDA).
172
173 # remember we have x as input features and y as output features
174
175 # parameters
176 degree = 4 # polynomial degree
177 lambdas = np.logspace(-4, 2, 100) # lambdas for ridge
178 kf = KFold(n_splits=10, shuffle=True, random_state=1)
179
180 # map input features into its polynomial space
181 poly = PolynomialFeatures(degree=degree)
182 x_poly = poly.fit_transform(x)
183
184 # store cross-validation results
185 mse_values = []
186
187 # perform cross-validation for each lambda
188 for alpha in lambdas:
189     ridge = Ridge(alpha=alpha)
190
191     # calc mse using cross_val_score
192     mse = -cross_val_score(ridge, x_poly, y, cv=kf, scoring='
        neg_mean_squared_error').mean()
193     mse_values.append(mse)
194
195 # find lambda with best minimum mse
196 best_lambda = lambdas[np.argmin(mse_values)]
197 print(f"Best lambda (w/ minimum MSE): {best_lambda}")
198
199 # fit model with best lambda
200 ridge_best = Ridge(alpha=best_lambda)
201 ridge_best.fit(x_poly, y)
202
203 # generate testing x values
204 x_fit = np.linspace(x.min(), x.max(), 1000).reshape(-1,1)
205 x_poly_fit = poly.transform(x_fit) # send x_fit to polynomial
    space

```

```

206 y_fit = ridge_best.predict(x_poly_fit)
207
208 # plot mse vs. ln(lambda)
209 plt.figure(figsize=(8,6))
210 plt.plot(np.log(lambdas), mse_values, label='MSE')
211 plt.xlabel('$\\log_e (\\lambda)$')
212 plt.ylabel('Cross-Validated MSE')
213 plt.title('Cross-Validated MSE v. $\\log_e (\\lambda)$')
214 plt.grid(True)
215 plt.legend()
216 plt.savefig('crossvalmse_vs_loglambda.png', dpi=300)
217 plt.close()

```

```

Degree: 1
MSE Training: 19458.420066787396
MSE Testing: 19657.90065574125
Degree: 2
MSE Training: 12843.093582390631
MSE Testing: 12953.786640428427
Degree: 3
MSE Training: 5.011946346822575
MSE Testing: 5.000928251849794
Degree: 4
MSE Training: 5.011930151180137
MSE Testing: 5.000749381529411
Degree: 5
MSE Training: 5.011819378630573
MSE Testing: 5.000656140106037
Best Polynomial Degree: 3
Best lambda (w/ minimum MSE): 0.24770763559917114

```

Figure 1: q1.py Terminal Output




Figure 2: Polynomial Model Degree 1




Figure 3: Polynomial Model Degree 2




Figure 4: Polynomial Model Degree 3




Figure 5: Polynomial Model Degree 4

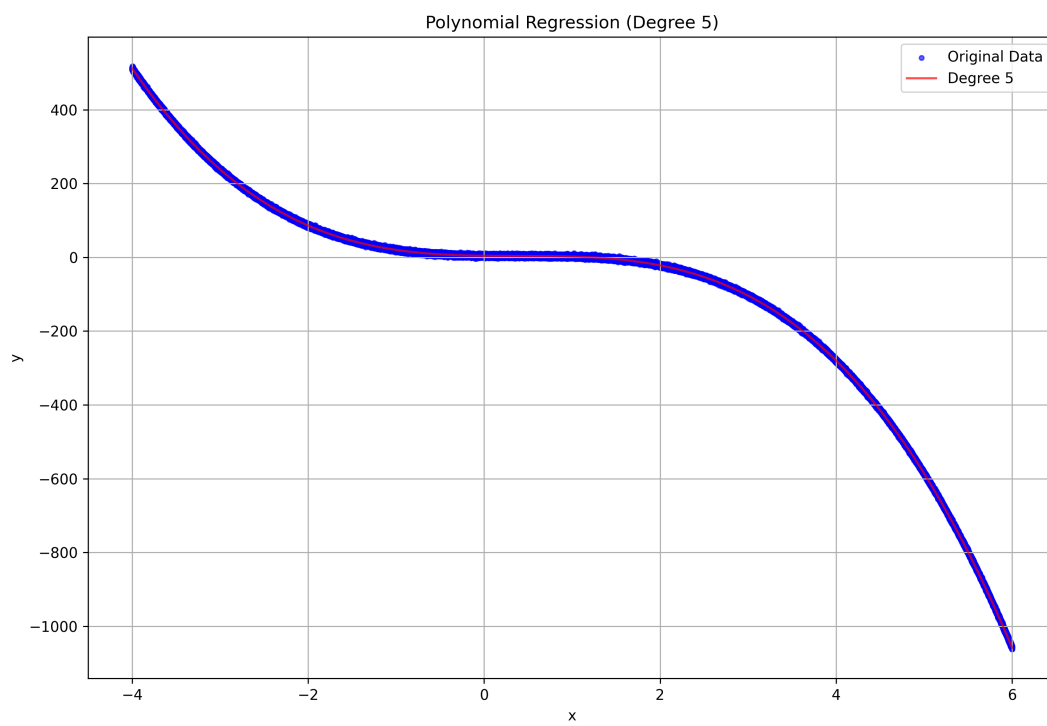


Figure 6: Polynomial Model Degree 5

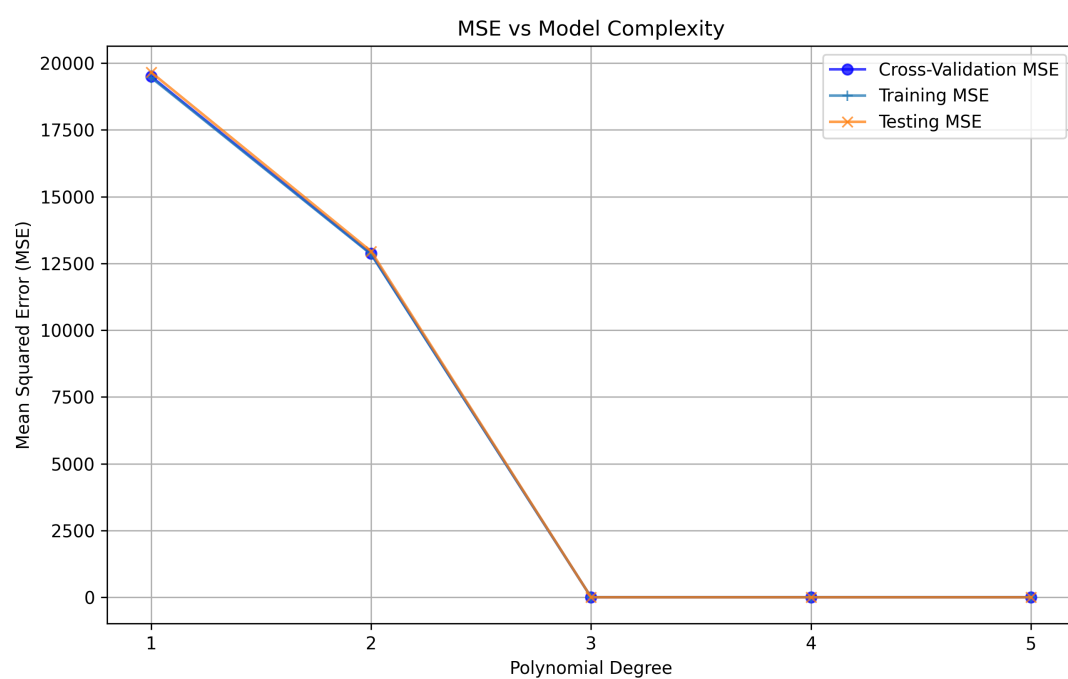


Figure 7: MSE v. Model Complexity

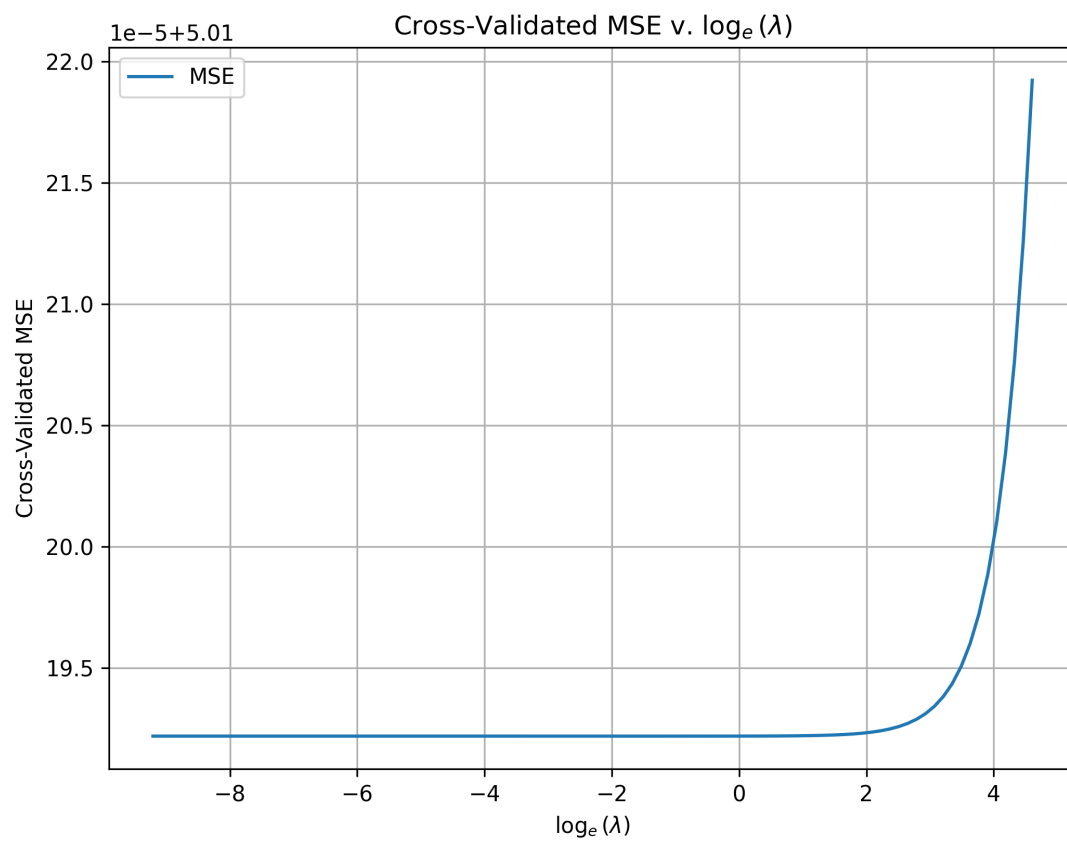


Figure 8: Cross-Validated MSE v. $\log_e \lambda$

Question 2

Solution.

q2.py

```
1 #####
2 # Chase Lotito - SIUC F24 #
3 # ECE469 - Intro to ML    #
4 # HW2 - Question 2        #
5 #####
6
7 # IMPORT LIBRARIES
8 import numpy as np
9 import pandas as pd
10 from sklearn.preprocessing import OrdinalEncoder    # For encoding
    categorical features
11 from sklearn.impute import SimpleImputer           # For adding
    missing values
12 from sklearn.preprocessing import StandardScaler   # For
    standardizing data
13 from sklearn.preprocessing import PolynomialFeatures
14 from sklearn.model_selection import cross_val_score,
    train_test_split, KFold
15 from sklearn.linear_model import LinearRegression
16 from sklearn.linear_model import Ridge
17 from sklearn.metrics import mean_squared_error
18 import matplotlib.pyplot as plt
19
20
21 # (A) DOWNLOAD HOUSING.CSV
22
23 # Get housing data
24 RAW_DATA = 'https://github.com/ageron/data/raw/main/housing/housing.
    csv'
25 housing = pd.read_csv(RAW_DATA)
26
27
28
29 # (B) DATA-PREPROCESSING FROM HW1
30
31 # Choose input features and output features (saved into numpy.
    ndarray type)
32 X = housing[
33     ['longitude',
34     'latitude',
35     'housing_median_age',
36     'total_rooms',
```

```

37         'total_bedrooms',
38         'population',
39         'households',
40         'median_income',
41         'ocean_proximity']
42     ].values
43 Y = housing[['median_house_value']].values
44
45 # Ocean Proximity is a categorical feature. Drop it or transform
    into numerical values (encode).
46
47 # Isolate the ocean_proximity data in input data X
48 ocean_proximity = X[:,8].reshape(-1,1)    # reshape(-1,1) to make 2D
    array for Ordinal
49 # Initialize the ordinal encoder
50 ordinal_encoder = OrdinalEncoder()
51 # Encode the ocean_proximity strings into numerical data
52 encoded_ocean = ordinal_encoder.fit_transform(ocean_proximity)
53 # Put the encoded version of ocean_proximity into input data X
54 X[:,8] = encoded_ocean.flatten()          # flatten to add 1D version
    of array back into X
55
56 # Clean the data by either dropping or replacing missing values
57
58 # Initialized SimpleImputer, will use the median to add missing
    entries
59 simple_imputer = SimpleImputer(strategy='median')
60
61 # Change X np ndarray into a Pandas Dataframe to use SimpleImputer
62 dX = pd.DataFrame(X)
63 dY = pd.DataFrame(Y)
64
65 # Perform SimpleImputer transformation, for both inputs X and
    outputs Y
66 imputed_data = simple_imputer.fit_transform(dX)
67 X = imputed_data
68 imputed_data = simple_imputer.fit_transform(dY)
69 Y = imputed_data
70
71 # Carry out feature scaling either via normalization or
    standardization.
72 std_scaler = StandardScaler()
73 scaled_data = std_scaler.fit_transform(X)
74 X = scaled_data
75 scaled_data = std_scaler.fit_transform(Y)
76 Y = scaled_data
77
78 # set test set size

```

```

79 test_size = 0.2
80
81 # split into testing and training set (both outputted as pd.
    DataFrames)
82 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=
    test_size, random_state=42)
83
84
85
86
87 # (C) USE LINEAR REGRESSION TO DEVELOP AN ML MODEL FOR PREDICTION OF
88 #     'MEDIAN_HOUSE_VALUE' FOR FUTURE INPUTS AND ANALYZE TEST ERRORS
89 #     EXPLICITLY EXPRESS THE CORRESPONDING OPTIMAL WEIGHTS AND THE
90 #     FINAL LEARNED MODEL. USE GRAPHICAL REPRESENTATIONS.
91
92 # initialize and train linear model
93 linear_regressor = LinearRegression()
94 linear_regressor.fit(X_train, Y_train)    # <-- train model here
95
96 # extract the optimal weights from the ML model
97 weights = linear_regressor.coef_.flatten()    # flatten makes it a
    normal list
98
99 # test model
100 Y_train_predicted = linear_regressor.predict(X_train)
101 Y_test_predicted = linear_regressor.predict(X_test)
102
103 # calculate mean square error
104 mse_train = mean_squared_error(Y_train, Y_train_predicted)
105 mse_test = mean_squared_error(Y_test, Y_test_predicted)
106
107 # VISUALIZATION
108 # print out results of linear model
109 print("-----")
110 print("LINEAR MODEL RESULTS")
111 print("-----")
112 print(f"Model Weights: {weights}")
113 print(f"MSE (train): {mse_train*100:.2f}%")
114 print(f"MSE (test): {mse_test*100:.2f}%")
115
116 # Plot predicted vs actual values
117 plt.scatter(Y_test, Y_test_predicted, marker='o', s=0.75, c='#32a852
    ', alpha=0.95, label='Pred v. Actual')    # plot predicted
    against actual, if diagonalized, well fit.
118 plt.plot([min(Y_test), max(Y_test)], [min(Y_test), max(Y_test)],
    label='Ideal Model', color='red', linewidth=2)
119 plt.xlabel("Actual Median House Value")
120 plt.ylabel("Predicted Median House Value")

```



```

121 plt.title("Predicted Median House Value vs Actual Median House Value
122 ")
123 plt.legend()
124 plt.savefig("plots\\q2_pred_vs_actual.png", dpi=300)
125 plt.close()
126 # Plot the Learned Weights
127 # Coefficients of the model
128 #feature_names = ['longitude', 'latitude', 'housing_median_age', '
129     total_rooms',
130     'total_bedrooms', 'population', 'households', '
131     median_income', 'ocean_proximity']
132 #plt.barh(feature_names, weights)
133 #plt.xlabel("Model Weights $w_i$")
134 #plt.title("Linear Regression Feature Weights")
135 #plt.show()
136 #plt.close()
137
138
139 # (D) USE CROSS-VALIDATION TECHNIQUES TO IMPROVE THE GENERALIZATION
140 # OF THE MODEL AND ANALYZE
141 # THE ROOT MEAN SQUARE ERROR (RMSE). USE GRAPHICAL ILLUSTRATIONS
142 #
143 # parameters
144 lambdas = np.logspace(-4, 2, 100) # lambdas for ridge
145 kf = KFold(n_splits=10, shuffle=True, random_state=1)
146 # store cross-validation results
147 mse_values = []
148
149 # perform cross-validation for each lambda
150 for alpha in lambdas:
151     ridge = Ridge(alpha=alpha)
152
153     # calc mse using cross_val_score
154     mse = -cross_val_score(ridge, X, Y, cv=kf, scoring='
155         neg_mean_squared_error').mean()
156     mse_values.append(mse)
157
158 # find lambda with best minimum mse
159 best_lambda = lambdas[np.argmin(mse_values)]
160 print("-----")
161 print("RIDGE REGULARIZATION")
162 print("-----")
163 print(f"Best lambda (w/ minimum MSE): {best_lambda:.2f}")

```

```

163 print(f"Minimum MSE: {np.min(mse_values)*100:.2f}%")
164
165 # fit model with best lambda
166 ridge_best = Ridge(alpha=best_lambda)
167 ridge_best.fit(X_train, Y_train)
168
169 # testing x values
170 Y_test_ridge_pred = ridge_best.predict(X_test)
171
172 # plot rmse vs. ln(lambda)
173 plt.figure(figsize=(8,6))
174 plt.plot(np.log(lambdas), np.sqrt(mse_values), c="#570710",label='
    MSE')
175 plt.xlabel('$\\log_e (\\lambda)$')
176 plt.ylabel('Cross-Validated RMSE')
177 plt.title('Q2: Cross-Validated RMSE v. $\\log_e (\\lambda)$')
178 plt.grid(True)
179 plt.legend()
180 plt.savefig('plots\\q2_crossvalmse_vs_loglambda.png', dpi=300)
181 plt.close()

```

We can plot, like in Figure 10, the predicted values against the actual values to see how well the learned model works. If the predicted values (green) lie along the same line as the actual values (red), then the model is ideal. However, we see the linear model has a large spread, which means the model is most likely not complex enough to model the input features.

```

-----
LINEAR MODEL RESULTS
-----
Model Weights: [-0.72960647 -0.77646856  0.12237936 -0.14310869  0.3233905  -0.3668318
 0.22907347  0.65294944  0.00340332]
MSE (train): 36.25%
MSE (test): 37.30%
-----
RIDGE REGULARIZATION
-----
Best lambda (w/ minimum MSE): 14.17
Minimum MSE: 36.55%

```

Figure 9: q2.py Terminal Output

From Figure 9, we can write the equation for the final learned model:

$$\begin{aligned}
 y(\mathbf{x}) = & -0.7296x_1 + -0.7765x_2 + 0.1224x_3 + -0.1431x_4 + 0.3234x_5 \\
 & + -0.3668x_6 + 0.2291x_7 + 0.6529x_8 + 0.0034x_9
 \end{aligned}$$

Where the input vector \mathbf{x} is defined as:

$$\mathbf{x} = \begin{bmatrix} x_1 : & \text{longitude} \\ x_2 : & \text{latitude} \\ x_3 : & \text{housing median age} \\ x_4 : & \text{total rooms} \\ x_5 : & \text{total bedrooms} \\ x_6 : & \text{population} \\ x_7 : & \text{households} \\ x_8 : & \text{median income} \\ x_9 : & \text{ocean proximity} \end{bmatrix}$$




Figure 10: Predicted Median House Value v. Actual Median House Value

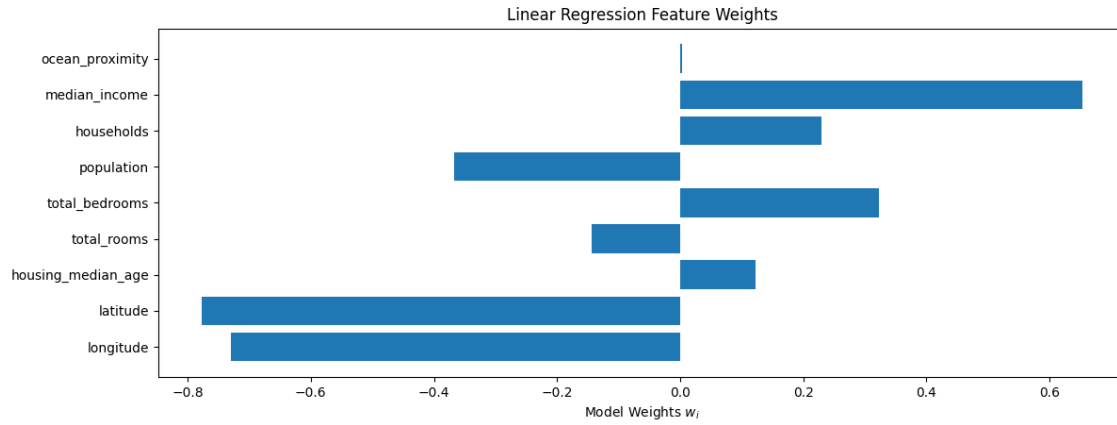


Figure 11: Model Weights

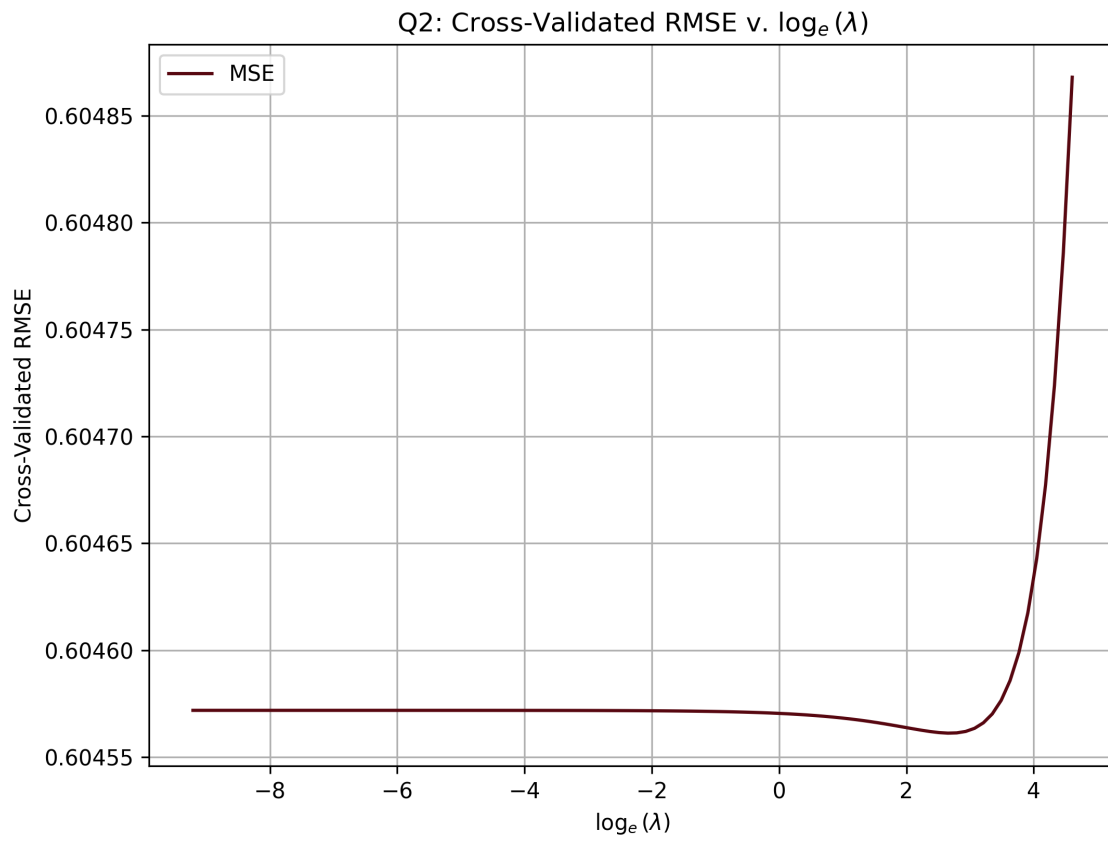


Figure 12: Cross-Validated RMSE v. $\log_e \lambda$