

ECE 469/ECE 568 Machine Learning

Textbook:

Machine Learning: a Probabilistic Perspective by Kevin Patrick Murphy

Southern Illinois University

September 13, 2024

A recap for the last lecture

- In the last lecture, we discussed the followings:
 - Introduction to regularization to mitigate overfitting
 - Types of regularization techniques for ML
 - Intuitions behind mathematical formulations for regularizers

Hyperparameters in machine learning models

- Recall that we learned about model parameters and hyperparameters:
 - Model training \rightarrow model parameters
 - Hyperparameter tuning \rightarrow best hyperparameters
- The behavior of the learning algorithm can be controlled by using several settings, which are termed "hyperparameters".
- These hyperparameters are not typically learned by the learning algorithm itself.

Hyperparameters in machine learning models

- Examples for hyperparameters:
 - ① In our polynomial regression example, there is a single hyperparameter: the degree of the polynomial (M) \rightarrow This acts as a "capacity" hyperparameter.
 - ② You have also learned that in regularization, the value of λ can be used to control the strength of weight decay. Hence, λ is another example of a hyperparameter.
- Sometimes, we may choose a learning architecture/framework setting to be a hyperparameter (that the learning algorithm does not learn) because this particular setting is difficult to optimize.
- Typically, we do not try to learn hyperparameters because it is not appropriate to learn that hyperparameter on the training set.

Validation set

- Recall that a test set consists examples coming from the same distribution as the training set. But data in test set and training set are independent.
- Note that the test set is not used in any way to learn the model, including its hyperparameters.
- A validation set can be used to choose appropriate values for hyperparameters.
- Data from the test set may not be used in the validation set.

Validation set

Validation set

- We may split the training data into two disjoint subsets: -> One subset is used to learn the model parameters: -> Training set.
- While the other subset is the validation set, which is used to estimate the generalization error during or after training, allowing for the hyperparameters to be updated accordingly.
- The subset of data used to guide the selection of hyperparameters is called the validation set.

Validation set

- Validation set can be used to tune/fine-tune hyperparameters.
- Validation data can be dedicated or shared with training data.
- For example, the training data can be split into a 80% training set and a 20% validation set.
- In cross-validation techniques, training and validation may be performed on the same set of training data.
- The training set is used to train the machine learning model, while the validation set is used to select the hyperparameters for the final model.
- The test set is typically independent from both training and validation data.
- Test data-set is used to estimate the final performance metrics such as the test error after all hyperparameter selections have been completed.

Training, validation and testing

- Recall the splitting of data-sets into training, validation, and test sets.

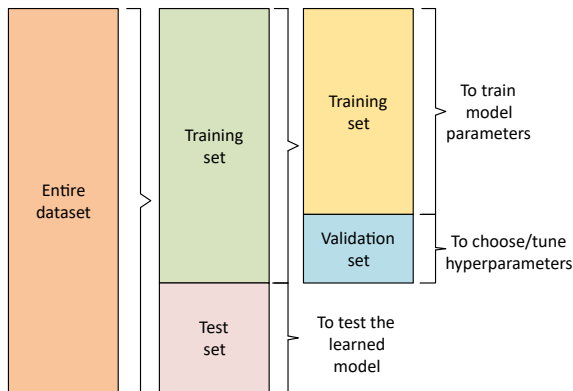


Figure: Splitting a data-set into training, validating, and testing sets

Regularization in machine learning

- Here, we formally define regularized loss function.

$$\underbrace{\tilde{J}(\mathbf{w})}_{\text{regularized loss}} = \underbrace{J(\mathbf{w})}_{\text{original loss function}} + \underbrace{\lambda}_{\text{regularization hyperparameter}} \times \underbrace{R(\mathbf{w})}_{\text{regularizer}}$$

- The balance between the regularized loss and original loss function is controlled by the regularization parameter λ .
- The most commonly used regularization is the l_2 regularization.

$$R_{\mathbf{w}} = \frac{1}{N} \|\mathbf{w}\|_2^2 = \frac{1}{N} \mathbf{w}^T \mathbf{w},$$

where N is the size of the data-set.

- $\|\mathbf{w}\|_2$ is l_2 -norm defined as $\|\mathbf{w}\| = \sqrt{\mathbf{w}^T \mathbf{w}} = \sqrt{w_0^2 + w_1^2 + \dots + w_n^2}$.
- It encourages the optimizer to find a model with the smallest l_2 norm.
- In deep learning, the most commonly used regularizer is l_2 regularization (a.k.a. weight-decay regularizer).

Gradient descent with l_2 regularization

- Recall that the Ridge (l_2) regularized loss is

$$\tilde{J}(\mathbf{w}) = J(\mathbf{w}) + \frac{\lambda}{N} \mathbf{w}^T \mathbf{w}$$

- The weight update in gradient descent with l_2 regularization is

$$\mathbf{w} := \mathbf{w} - \alpha \left(\nabla \tilde{J}(\mathbf{w}) \right)$$

- The gradient of the regularized loss is given by

$$\nabla \tilde{J}(\mathbf{w}) = \nabla J(\mathbf{w}) + \frac{2\lambda}{N} \mathbf{w}$$

- Then, the weight update is given by

$$\mathbf{w} := \mathbf{w} - \alpha \left(\nabla J(\mathbf{w}) + \frac{2\lambda}{N} \mathbf{w} \right) = \left(1 - \frac{2\alpha\lambda}{N} \right) \mathbf{w} - \alpha \nabla J(\mathbf{w})$$

- The term $(1 - 2\alpha\lambda/N)$ in-front of \mathbf{w} corresponds to a weight decay.
- This is because the gradient descent with learning rate α on the regularized loss $\tilde{J}(\mathbf{w})$ is now equivalent to decaying \mathbf{w} by a scalar factor of $(1 - \alpha\lambda)$ and then applying the standard gradient descent.

Other regularizers

- Another regularizer can be chosen as $R(\mathbf{w}) = \|\mathbf{w}\|_1/N = \frac{1}{N} \sum_{i=1}^M |w_i|$, where N is the size of the training data-set.
- This is termed as the least absolute shrinkage and selection operator (LASSO).
- In deep learning, other common regularizers include dropout, data augmentation, regularizing the spectral norm of the weight matrices, etc.
- Regularization in deep learning is an active research area.

Explicit vs. implicit regularization

- In explicit regularization, we explicitly add a term to the optimization objective function.
- Implicit regularization is all other forms of regularization such as early stopping, using a robust loss function, and discarding outliers.

Analytical solution for the optimal weights with Ridge regression

- Let us consider a linear regression model: $f(\mathbf{x}, \mathbf{w}) = \mathbf{w}^T \mathbf{x}$.
- The modified loss function with Ridge regularizer is given by

$$J(\mathbf{w}) = \frac{1}{N} \sum_{m=1}^M \left[y^{(m)} - \mathbf{w}^T \mathbf{x}^{(m)} \right]^2 + \frac{\lambda}{N} \sum_{i=0}^n w_i^2$$

- This loss function can also be written as

$$J(\mathbf{w}) = \frac{1}{N} (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) + \frac{\lambda}{N} \mathbf{w}^T \mathbf{w}$$

where $\mathbf{X} = [\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}]^T$ and $\mathbf{y} = [y^{(1)}, y^{(2)}, \dots, y^{(N)}]^T$

Analytical solution for Ridge regression

- This loss function can be expanded as

$$\begin{aligned}J(\mathbf{w}) &= \frac{1}{N} (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) + \frac{\lambda}{N} \mathbf{w}^T \mathbf{w} \\&= \frac{1}{N} ((\mathbf{X}\mathbf{w})^T - \mathbf{y}^T) (\mathbf{X}\mathbf{w} - \mathbf{y}) + \frac{\lambda}{N} \mathbf{w}^T \mathbf{w} \\&= \frac{1}{N} (\mathbf{w}^T \mathbf{X}^T - \mathbf{y}^T) (\mathbf{X}\mathbf{w} - \mathbf{y}) + \frac{\lambda}{N} \mathbf{w}^T \mathbf{w} \\&= \frac{1}{N} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}) + \frac{\lambda}{N} \mathbf{w}^T \mathbf{w}\end{aligned}$$

- Useful identities in multivariate calculus:

$f(\mathbf{w})$	$\frac{\partial f}{\partial \mathbf{w}}$
$\mathbf{w}^T \mathbf{x}$	\mathbf{x}
$\mathbf{x}^T \mathbf{w}$	\mathbf{x}
$\mathbf{w}^T \mathbf{w}$	$2\mathbf{w}$
$\mathbf{w}^T \mathbf{C} \mathbf{w}$	$2\mathbf{C} \mathbf{w}$

Analytical solution for Ridge regression

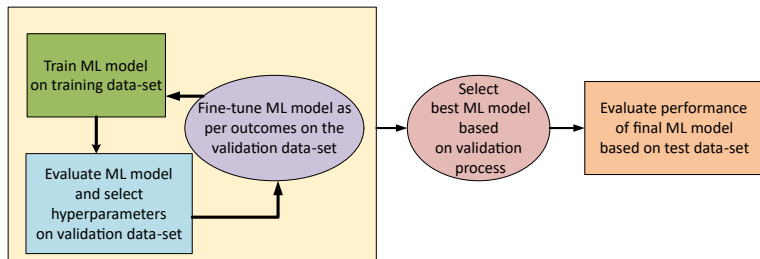
- The gradient vector of this loss function can be computed as

$$\begin{aligned}\nabla_{\mathbf{w}} J(\mathbf{w}) &= \begin{bmatrix} \frac{\partial J(\mathbf{w})}{\partial w_0} \\ \frac{\partial J(\mathbf{w})}{\partial w_1} \\ \vdots \\ \frac{\partial J(\mathbf{w})}{\partial w_n} \end{bmatrix} = \frac{1}{N} \nabla_{\mathbf{w}} \left(\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2 \mathbf{w}^T \mathbf{X}^T \mathbf{y} + \mathbf{y}^T \mathbf{y} + \frac{\lambda}{N} \mathbf{w}^T \mathbf{w} \right) \\ &= \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{y} + \frac{2\lambda}{N} \mathbf{w}\end{aligned}$$

- Next, the optimal weight vector that minimizes the loss function $J(\mathbf{w})$ can be computed as

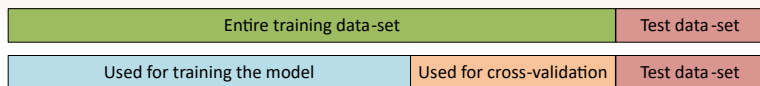
$$\begin{aligned}\nabla_{\mathbf{w}} J(\mathbf{w}) &= \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{y} + \frac{2\lambda}{N} \mathbf{w} = 0 \\ \mathbf{w}^* &= (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}\end{aligned}$$

Validation: Model selection via cross-validation



- Next, we are going to discuss validation techniques to finalize hyperparameters and choosing the final model.

Hold-out (simple) cross-validation



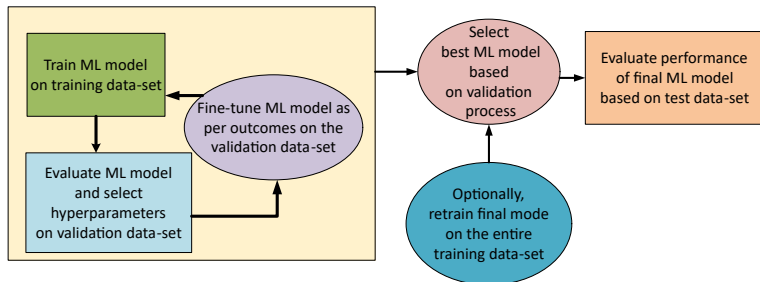
- Let us assume that we have a finite set of models denoted by

$$\mathcal{M} = \{M_1, \dots, M_D\},$$

where model M_i is the i th degree polynomial regression model.


- The hold-out (simple) cross-validation works as follows:
 - Randomly split the training data set \mathcal{S} into two sets, namely $\mathcal{S}_{\text{train}}$ (say, 80% of the training data) and \mathcal{S}_{cv} (the remaining 20%), where \mathcal{S}_{cv} is called the hold-out cross-validation set.
 - Train each model M_i on $\mathcal{S}_{\text{train}}$ only, to find a set of hypothesis h_i .
 - Select and output the hypothesis h_i that results in the smallest error $E_{\mathcal{S}_{\text{cv}}}(h_i)$ using the hold-out cross-validation set \mathcal{S}_{cv} .
- Note that $E_{\mathcal{S}_{\text{cv}}}(h_i)$ denotes the mean square error (MSE) of h_i on the cross-validation set \mathcal{S}_{cv} , and it is explicitly referred to as the validation error.

Model selection via cross-validation



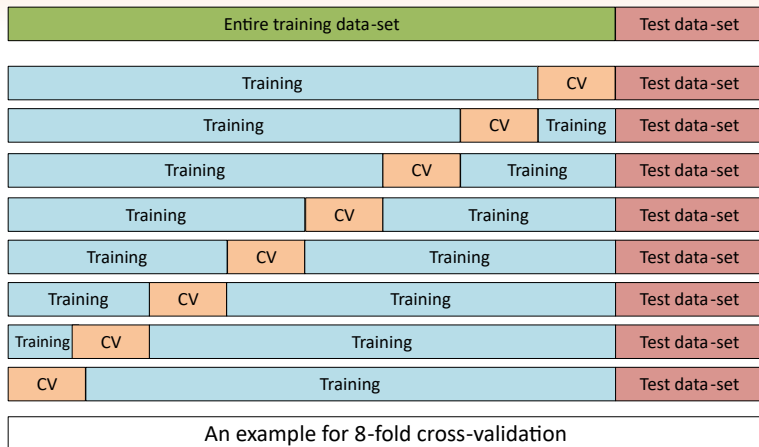
- Optionally, the above third step may also be replaced with selecting the model M_i according to $\arg \min_i E_{\mathcal{S}_{cv}}(h_i)$, and then retraining M_i on the entire training set \mathcal{S} .

Cons of hold-out (simple) cross-validation and remedies



- Cross-validation techniques should allow us to make more efficient use of the data that is available to us.
- One disadvantage of using hold-out cross-validation is that it wastes a pre-defined portion of (say 20% of the training data).
- This is crucial in learning problems in which data is scarce and costly.
- The remedy is to use k -fold cross-validation, which holds out less data each time.

k -fold cross-validation



- k -fold cross-validation involves randomly dividing the data-set into k -folds/groups of approximately equal size.
- The first fold is treated as a validation set, and the model is fit on the remaining $k - 1$ folds.

k -fold cross validation

- The mean square error (MSE) evaluated on the first held-out fold is denoted by MSE_1 .
- This process is repeated k times, each time, a different fold is treated as a validation set.
- For each hold-out fold, the MSE is calculated as $\text{MSE}_1, \text{MSE}_2, \dots, \text{MSE}_k$.
- The final k -fold CV MSE is computed by averaging these K errors as

$$\text{MSE}_{\text{CV}} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$

k -fold cross-validation

- k -fold cross-validation algorithm is given below:

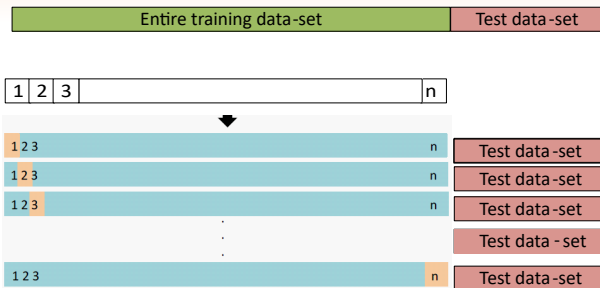
An algorithm for k -fold cross-validation

- ➊ Randomly split the training set \mathcal{S} into k disjoint subsets of N/k training examples each, namely $\{\mathcal{S}_1, \dots, \mathcal{S}_k\}$. Here $N = |\mathcal{S}|$ is the size of the training set.
- ➋ For each model M_i , evaluate the followings:
 - For $j = 1, \dots, k$, first train the model M_i on $\mathcal{S}_1 \cup \mathcal{S}_2 \cup \dots \mathcal{S}_{j-1} \cup \mathcal{S}_{j+1} \cup \dots \cup \mathcal{S}_k$ (i.e., train on all the data except \mathcal{S}_j) to obtain some hypothesis h_{ij} .
 - Then, test the hypothesis h_{ij} on \mathcal{S}_j , to compute $E_{\mathcal{S}_j}(h_{ij})$.
 - Finally, the estimated cross-validation error of model M_i is calculated as the average of the $E_{\mathcal{S}_j}(h_{ij})$ (averaged over j).
- ➌ Pick the model M_i with the lowest estimated cross-validation error, and retrain that model on the entire training set \mathcal{S} .

k -fold cross-validation

- The resulting hypothesis is then output as the final learning model.
- An independent test data-set is used to evaluate the performance of this final model.
- A typical choice for the number of folds is $k = 10$.
- k -fold cross validation is computationally more expensive than hold-out cross-validation, since it needs train to each model k times.

Leave-one-out cross-validation (LOOCV)



- When data is really scarce, the extreme choice of $k = n$, the size of the training data-set, can be considered to leave out as little data as possible each time.
- Since we are going to hold out one training example at a time, this method is called leave-one-out cross-validation.
- In this setting, we repeatedly train on all but one of the training examples in \mathcal{S} , and validate on that held-out example.

Leave-one-out cross-validation (LOOCV)

- Then, the resulting $n = k$ errors can be averaged together to obtain the estimate of the cross-validation error for a learning model.
- The final LOOCV mean square error is computed by averaging all n errors as

$$\text{MSE}_{\text{CV}} = \frac{1}{n} \sum_{i=1}^n \text{MSE}_i$$

- LOOCV has a two of key advantages over the simple hold-out validation set technique.
- First, LOOCV has far less bias because it repeatedly fits the model on training sets that contain $n - 1$ examples, which is almost as many as are in the entire training data set.
- Consequently, LOOCV would not to overestimate the validation error as much as the hold-out simple validation set technique does.